

CNN's on MNIST

```
In [0]: # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist\_cnn.py

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                  activation='relu',
                  input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
```

```

optimizer=keras.optimizers.Adadelta(),
metrics=['accuracy'])

model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

```

In [1]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensor
flow" use this command
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense,Activation
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout,Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

```

Using TensorFlow backend.

```

In [2]: # the data, shuffled and split between train and test sets
(X_train, y_train), (x_test, y_test) = mnist.load_data()

```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```

In [0]: # some model parameters
input_dim = X_train.shape[1]
num_clas = 10
batch_size = 128
no_epoch = 20
img_row, img_col = 28,28

```

```

In [0]: if K.image_data_format() == 'channels_first':
        x_train = x_train.reshape(x_train.shape[0], 1, img_row, img_col)
        x_test = x_test.reshape(x_test.shape[0], 1, img_row, img_col)
        input_shape = (1, img_row, img_col)
    else:
        X_train = X_train.reshape(X_train.shape[0], img_row, img_col, 1)
        x_test = x_test.reshape(x_test.shape[0], img_row, img_col, 1)
        input_shape = (img_row, img_col, 1)

```

```

In [5]: print("Number of training examples :", X_train.shape[0], X_train.shape[1],
X_train.shape[2])

```

Number of training examples : 60000 28 28

```
In [6]: print("Number of training examples :", X_train.shape[0], "and each image is
of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", x_test.shape[0], "and each image is
of shape (%d, %d)"%(x_test.shape[1], x_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
 Number of training examples : 10000 and each image is of shape (28, 28)

```
In [7]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is
of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", x_test.shape[0], "and each image is
of shape (%d)"%(x_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (28)
 Number of training examples : 10000 and each image is of shape (28)

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms Lets try to normalize
the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
x_test = x_test/255
```

```
In [9]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0,
0, 0]
# this conversion needed for MLPs

y_train = keras.utils.to_categorical(y_train, num_clas)
y_test = keras.utils.to_categorical(y_test, num_clas)

print("After converting the output into a vector : ", type(y_train))
```

Class label of first image : 5
 After converting the output into a vector : <class 'numpy.ndarray'>

```
In [10]: print(y_train.shape)
print(y_test.shape)
```

(60000, 10)
 (10000, 10)

3-layer ConvNet + Relu + adam

No Batch normalization and Dropout

```
In [0]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_md1_res(x_val, trn_los, tst_los, tst_scr, tst_acc):
    # Visualize loss history
    plt.figure(figsize=(16,16))
    plt.plot(x_val, trn_los, 'r--')
    plt.plot(x_val, tst_los, 'b-')
    plt.legend(['Training Loss', 'Test Loss'])
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.grid()
    plt.show();

    print('Test score:', tst_scr)
    print('Test accuracy:', tst_acc)
```

```
In [0]: #ConvNet
mdl_1 = Sequential()
mdl_1.add(Conv2D(32, kernel_size=(7,7), activation='relu', input_shape=(input_shape)))
mdl_1.add(Conv2D(64, kernel_size=(5,5), activation='relu'))
mdl_1.add(MaxPooling2D(pool_size=(2,2)))
mdl_1.add(Conv2D(128, kernel_size=(5,5), activation='relu'))
mdl_1.add(Flatten())
mdl_1.add(Dense(128, activation='relu'))
mdl_1.add(Dense(num_clas, activation='softmax'))
mdl_1.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
=====		
conv2d_28 (Conv2D)	(None, 22, 22, 32)	1600
conv2d_29 (Conv2D)	(None, 18, 18, 64)	51264
max_pooling2d_10 (MaxPooling)	(None, 9, 9, 64)	0
conv2d_30 (Conv2D)	(None, 5, 5, 128)	204928
flatten_10 (Flatten)	(None, 3200)	0
dense_19 (Dense)	(None, 128)	409728
dense_20 (Dense)	(None, 10)	1290
=====		
Total params: 668,810		
Trainable params: 668,810		
Non-trainable params: 0		

```
In [0]: #set optimizer and loss
mdl_1.compile(optimizer=keras.optimizers.Adam(),loss=keras.losses.categorical_crossentropy, metrics=['accuracy'])
```

```
In [0]: history = mdl_1.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch,v  
erbose=1,validation_data=(x_test,y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 10s 166us/step - loss: 0.1582 - acc: 0.9500 - val_loss: 0.0419 - val_acc: 0.9869

Epoch 2/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0429 - acc: 0.9873 - val_loss: 0.0362 - val_acc: 0.9878

Epoch 3/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0308 - acc: 0.9905 - val_loss: 0.0287 - val_acc: 0.9912

Epoch 4/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0229 - acc: 0.9925 - val_loss: 0.0274 - val_acc: 0.9917

Epoch 5/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0183 - acc: 0.9941 - val_loss: 0.0350 - val_acc: 0.9890

Epoch 6/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0155 - acc: 0.9948 - val_loss: 0.0275 - val_acc: 0.9912

Epoch 7/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0132 - acc: 0.9960 - val_loss: 0.0319 - val_acc: 0.9918

Epoch 8/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0111 - acc: 0.9966 - val_loss: 0.0240 - val_acc: 0.9933

Epoch 9/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0119 - acc: 0.9960 - val_loss: 0.0246 - val_acc: 0.9921

Epoch 10/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0100 - acc: 0.9968 - val_loss: 0.0327 - val_acc: 0.9914

Epoch 11/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0078 - acc: 0.9973 - val_loss: 0.0421 - val_acc: 0.9907

Epoch 12/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0074 - acc: 0.9976 - val_loss: 0.0298 - val_acc: 0.9924

Epoch 13/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0072 - acc: 0.9977 - val_loss: 0.0282 - val_acc: 0.9938

Epoch 14/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0077 - acc: 0.9975 - val_loss: 0.0314 - val_acc: 0.9931

Epoch 15/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0070 - acc: 0.9978 - val_loss: 0.0277 - val_acc: 0.9929

Epoch 16/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0053 - acc: 0.9984 - val_loss: 0.0281 - val_acc: 0.9925

Epoch 17/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0038 - acc: 0.9989 - val_loss: 0.0406 - val_acc: 0.9920

Epoch 18/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0065 - acc: 0.9981 - val_loss: 0.0246 - val_acc: 0.9936

Epoch 19/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0053

- acc: 0.9983 - val_loss: 0.0295 - val_acc: 0.9933

Epoch 20/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0060


- acc: 0.9983 - val_loss: 0.0289 - val_acc: 0.9938

```
In [0]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history.history['loss']
test_loss = history.history['val_loss']

# Create count of the number of epochs
score = mdl_1.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.028870063201270432

Test accuracy: 0.9938

Hyper Parameter tuning Dropout rate for 3-Layer convNet

```
In [ ]: no_cvnets = 8
mdl = [0] * no_cvnets
history = [0] * no_cvnets
nme = ["Dropout=0.0", "Dropout=0.1", "Dropout=0.2", "Dropout=0.3", "Dropout=0.4", "Dropout=0.5", "Dropout=0.6", "Dropout=0.7"]

for i in range(no_cvnets):
    #ConvNet
    mdl[i] = Sequential()
    mdl[i].add(Conv2D(32, kernel_size=(7,7), activation='relu', input_shape=(input_shape)))
    mdl[i].add(Conv2D(64, kernel_size=(5,5), activation='relu'))
    mdl[i].add(MaxPooling2D(pool_size=(2,2)))
    mdl[i].add(Conv2D(128, kernel_size=(5,5), activation='relu'))
    mdl[i].add(Flatten())
    mdl[i].add(Dense(128, activation='relu'))
    mdl[i].add(Dropout(rate=(i*0.01)))
    mdl[i].add(Dense(num_clas, activation='softmax'))
    mdl[i].compile(optimizer="adam", loss="categorical_crossentropy", metrics=["accuracy"])

for j in range(no_cvnets):
    history[j] = mdl[j].fit(X_train, y_train, batch_size=batch_size, epochs=no_epoch,
        validation_data = (x_test, y_test), verbose=1)

for k in range(no_cvnets):
    print("CNN {0}: Epochs={1:d}, Train accuracy={2:.5f}, Validation accuracy={3:.5f}".format(
        nme[k], no_epoch, max(history[k].history['acc']), max(history[k].history['val_acc'])))
```

CNN Dropout=0.0: Epochs=20, Train accuracy=0.99913, Validation accuracy=0.99410
 CNN Dropout=0.1: Epochs=20, Train accuracy=0.99910, Validation accuracy=0.99360
 CNN Dropout=0.2: Epochs=20, Train accuracy=0.99882, Validation accuracy=0.99400
 CNN Dropout=0.3: Epochs=20, Train accuracy=0.99883, Validation accuracy=0.99290
 CNN Dropout=0.4: Epochs=20, Train accuracy=0.99878, Validation accuracy=0.99330
 CNN Dropout=0.5: Epochs=20, Train accuracy=0.99910, Validation accuracy=0.99320
 CNN Dropout=0.6: Epochs=20, Train accuracy=0.99915, Validation accuracy=0.99340
 CNN Dropout=0.7: Epochs=20, Train accuracy=0.99865, Validation accuracy=0.99340

3-Layer ConvNet + ReLu + Adam + Dropout

No Batch Normalization with Dropout rate 0.0

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20
```

```
In [0]: #ConvNet
mdl_2 = Sequential()
mdl_2.add(Conv2D(32,kernel_size=(7,7), activation='relu',input_shape=(input
_shape)))
mdl_2.add(Conv2D(64,kernel_size=(5,5), activation='relu'))
mdl_2.add(MaxPooling2D(pool_size=(2,2)))
mdl_2.add(Conv2D(128,kernel_size=(5,5), activation='relu'))
mdl_2.add(Dropout(rate=0.0))
mdl_2.add(Flatten())
mdl_2.add(Dense(128, activation='relu'))
mdl_2.add(Dropout(rate=0.0))
mdl_2.add(Dense(num_clas,activation='softmax'))
mdl_2.summary()
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
conv2d_31 (Conv2D)	(None, 22, 22, 32)	1600
conv2d_32 (Conv2D)	(None, 18, 18, 64)	51264
max_pooling2d_11 (MaxPooling)	(None, 9, 9, 64)	0
conv2d_33 (Conv2D)	(None, 5, 5, 128)	204928
dropout_10 (Dropout)	(None, 5, 5, 128)	0
flatten_11 (Flatten)	(None, 3200)	0
dense_21 (Dense)	(None, 128)	409728
dropout_11 (Dropout)	(None, 128)	0
dense_22 (Dense)	(None, 10)	1290
Total params: 668,810		
Trainable params: 668,810		
Non-trainable params: 0		

```
In [0]: #set optimizer and Loss
mdl_2.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['a
ccuracy'])
```

```
In [0]: history_2 = mdl_2.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch  
,verbose=1,validation_data=(x_test,y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 10s 170us/step - loss: 0.1612 - acc: 0.9502 - val_loss: 0.0455 - val_acc: 0.9856

Epoch 2/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0444 - acc: 0.9859 - val_loss: 0.0318 - val_acc: 0.9895

Epoch 3/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0297 - acc: 0.9906 - val_loss: 0.0287 - val_acc: 0.9913

Epoch 4/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0225 - acc: 0.9932 - val_loss: 0.0336 - val_acc: 0.9882

Epoch 5/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0198 - acc: 0.9938 - val_loss: 0.0451 - val_acc: 0.9870

Epoch 6/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0144 - acc: 0.9956 - val_loss: 0.0331 - val_acc: 0.9904

Epoch 7/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0134 - acc: 0.9959 - val_loss: 0.0363 - val_acc: 0.9888

Epoch 8/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0125 - acc: 0.9960 - val_loss: 0.0329 - val_acc: 0.9909

Epoch 9/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0095 - acc: 0.9969 - val_loss: 0.0394 - val_acc: 0.9899

Epoch 10/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0085 - acc: 0.9974 - val_loss: 0.0282 - val_acc: 0.9924

Epoch 11/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0088 - acc: 0.9971 - val_loss: 0.0329 - val_acc: 0.9912

Epoch 12/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0083 - acc: 0.9974 - val_loss: 0.0331 - val_acc: 0.9927

Epoch 13/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0077 - acc: 0.9974 - val_loss: 0.0390 - val_acc: 0.9910

Epoch 14/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0053 - acc: 0.9986 - val_loss: 0.0571 - val_acc: 0.9889

Epoch 15/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0073 - acc: 0.9978 - val_loss: 0.0289 - val_acc: 0.9934

Epoch 16/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0056 - acc: 0.9982 - val_loss: 0.0357 - val_acc: 0.9917

Epoch 17/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0066 - acc: 0.9979 - val_loss: 0.0357 - val_acc: 0.9929

Epoch 18/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0028 - acc: 0.9992 - val_loss: 0.0349 - val_acc: 0.9925

Epoch 19/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0042

- acc: 0.9987 - val_loss: 0.0338 - val_acc: 0.9924

Epoch 20/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0064


- acc: 0.9982 - val_loss: 0.0300 - val_acc: 0.9928

```
In [0]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_2.history['loss']
test_loss = history_2.history['val_loss']

# Create count of the number of epochs
score = mdl_2.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.030026011430384324

Test accuracy: 0.9928

5-Layer ConvNet + ReLu + Adam

Batch Normalization with No DropOut

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20

#ConvNet
mdl_3 = Sequential()
mdl_3.add(Conv2D(32,kernel_size=(3,3), activation='relu',input_shape=(input_shape)))
mdl_3.add(Conv2D(32,kernel_size=(3,3), activation='relu'))
mdl_3.add(MaxPooling2D(pool_size=(2,2)))
mdl_3.add(Conv2D(32,kernel_size=(5,5), activation='relu'))
mdl_3.add(BatchNormalization())
mdl_3.add(Conv2D(64,kernel_size=(3,3), activation='relu'))
mdl_3.add(MaxPooling2D(pool_size=(2,2)))
mdl_3.add(Conv2D(128,kernel_size=(2,2), activation='relu'))
mdl_3.add(BatchNormalization())
mdl_3.add(Flatten())
mdl_3.add(Dense(128, activation='relu'))
mdl_3.add(Dense(num_clas,activation='softmax'))
mdl_3.summary()

mdl_3.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_3 = mdl_3.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch,verbose=1,validation_data=(x_test,y_test))
```


WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:2041: The name tf.nn.fused_batch_norm is deprecated. Please use tf.compat.v1.nn.fused_batch_norm instead.

Model: "sequential_12"

Layer (type)	Output Shape	Param #
conv2d_34 (Conv2D)	(None, 26, 26, 32)	320
conv2d_35 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_12 (MaxPooling)	(None, 12, 12, 32)	0
conv2d_36 (Conv2D)	(None, 8, 8, 32)	25632
batch_normalization_1 (Batch Normalization)	(None, 8, 8, 32)	128
conv2d_37 (Conv2D)	(None, 6, 6, 64)	18496
max_pooling2d_13 (MaxPooling)	(None, 3, 3, 64)	0
conv2d_38 (Conv2D)	(None, 2, 2, 128)	32896
batch_normalization_2 (Batch Normalization)	(None, 2, 2, 128)	512
flatten_12 (Flatten)	(None, 512)	0
dense_23 (Dense)	(None, 128)	65664
dense_24 (Dense)	(None, 10)	1290
Total params: 154,186		
Trainable params: 153,866		
Non-trainable params: 320		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 184us/step - loss: 0.1177 - acc: 0.9638 - val_loss: 0.0785 - val_acc: 0.9763

Epoch 2/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0399 - acc: 0.9873 - val_loss: 0.0716 - val_acc: 0.9784

Epoch 3/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0281 - acc: 0.9908 - val_loss: 0.0401 - val_acc: 0.9883

Epoch 4/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0211 - acc: 0.9936 - val_loss: 0.0626 - val_acc: 0.9827

Epoch 5/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0185 - acc: 0.9940 - val_loss: 0.1290 - val_acc: 0.9661

Epoch 6/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0176 - acc: 0.9942 - val_loss: 0.0317 - val_acc: 0.9899

Epoch 7/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0136


```
- acc: 0.9953 - val_loss: 0.0361 - val_acc: 0.9890
Epoch 8/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0137
- acc: 0.9956 - val_loss: 0.0337 - val_acc: 0.9902
Epoch 9/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0103
- acc: 0.9967 - val_loss: 0.0438 - val_acc: 0.9891
Epoch 10/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0128
- acc: 0.9957 - val_loss: 0.0574 - val_acc: 0.9842
Epoch 11/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0097
- acc: 0.9965 - val_loss: 0.0413 - val_acc: 0.9882
Epoch 12/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0092
- acc: 0.9970 - val_loss: 0.0547 - val_acc: 0.9873
Epoch 13/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0087
- acc: 0.9974 - val_loss: 0.0710 - val_acc: 0.9851
Epoch 14/20
60000/60000 [=====] - 9s 150us/step - loss: 0.0094
- acc: 0.9968 - val_loss: 0.0365 - val_acc: 0.9916
Epoch 15/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0073
- acc: 0.9976 - val_loss: 0.0327 - val_acc: 0.9922
Epoch 16/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0051
- acc: 0.9983 - val_loss: 0.0328 - val_acc: 0.9922
Epoch 17/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0071
- acc: 0.9978 - val_loss: 0.0433 - val_acc: 0.9907
Epoch 18/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0080
- acc: 0.9973 - val_loss: 0.0372 - val_acc: 0.9915
Epoch 19/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0082
- acc: 0.9972 - val_loss: 0.0532 - val_acc: 0.9882
Epoch 20/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0041
- acc: 0.9986 - val_loss: 0.0294 - val_acc: 0.9934
```

```
In [0]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_3.history['loss']
test_loss = history_3.history['val_loss']

# Create count of the number of epochs
score = mdl_3.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.029422459162015413

Test accuracy: 0.9934

7-Layer ConvNet + Adam + Relu

Batch Normalization with Dropout rate 0.0

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20

#ConvNet
mdl_4 = Sequential()
mdl_4.add(Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=(input_shape)))
mdl_4.add(Conv2D(32, kernel_size=(3,3), activation='relu'))
mdl_4.add(Conv2D(32, kernel_size=(3,3), activation='relu'))
mdl_4.add(Dropout(rate=0.0))

mdl_4.add(Conv2D(64, kernel_size=(3,3), activation='relu'))
mdl_4.add(Conv2D(64, kernel_size=(3,3), activation='relu'))

mdl_4.add(Dropout(rate=0.0))
mdl_4.add(BatchNormalization())

mdl_4.add(Conv2D(64, kernel_size=(5,5), activation='relu'))
mdl_4.add(Conv2D(128, kernel_size=(2,2), activation='relu'))

mdl_4.add(BatchNormalization())
mdl_4.add(Flatten())
mdl_4.add(Dense(128, activation='relu'))
mdl_4.add(Dense(num_clas, activation='softmax'))

mdl_4.summary()

mdl_4.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history_4 = mdl_4.fit(X_train, y_train, batch_size=batch_size, epochs=no_epoch, verbose=1, validation_data=(x_test, y_test))
```

```


In [0]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_4.history['loss']
test_loss = history_4.history['val_loss']

# Create count of the number of epochs
score = mdl_4.evaluate(x_test, y_test, verbose=0)
plt_mdl_res(epoch_count, training_loss, test_loss, score[0], score[1])

```



Test score: 0.08437914842601021

Test accuracy: 0.981

Hyperparameter tuning Dropout rate for 7-Layer ConvNet

```

In [ ]: no_cnvnets = 8
mdl = [0] * no_cnvnets
history = [0] * no_cnvnets
nme = ["Dropout=0.0", "Dropout=0.1", "Dropout=0.2", "Dropout=0.3", "Dropout=0.4", "Dropout=0.5", "Dropout=0.6", "Dropout=0.7"]

for i in range(no_cnvnets):
    #ConvNet
    mdl[i] = Sequential()
    mdl[i].add(Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=(
input_shape)))
    mdl[i].add(Conv2D(32, kernel_size=(3,3), activation='relu'))
    mdl[i].add(Conv2D(32, kernel_size=(3,3), activation='relu'))
    mdl[i].add(Dropout(rate=(i*0.01)))

    mdl[i].add(Conv2D(64, kernel_size=(3,3), activation='relu'))
    mdl[i].add(Conv2D(64, kernel_size=(3,3), activation='relu'))

    mdl[i].add(Dropout(rate=(i*0.01)))
    mdl[i].add(BatchNormalization())

    mdl[i].add(Conv2D(64, kernel_size=(5,5), activation='relu'))
    mdl[i].add(Conv2D(128, kernel_size=(2,2), activation='relu'))

    mdl[i].add(BatchNormalization())
    mdl[i].add(Flatten())
    mdl[i].add(Dense(128, activation='relu'))
    mdl[i].add(Dense(num_clas, activation='softmax'))
    mdl[i].compile(optimizer="adam", loss="categorical_crossentropy", metri
cs=["accuracy"])

for j in range(no_cnvnets):
    history[j] = mdl[j].fit(X_train, y_train, batch_size=batch_size, epochs
= no_epoch,
    validation_data = (x_test, y_test), verbose=1)
    print('Processed model:', j)

for k in range(no_cnvnets):
    print("CNN {0}: Epochs={1:d}, Train accuracy={2:.5f}, Validation accura
cy={3:.5f}".format(
        nme[k], no_epoch, max(history[k].history['acc']), max(history[k].histo
ry['val_acc'])))

```

CNN Dropout=0.0: Epochs=20, Train accuracy=0.99800, Validation accuracy=0.99200
 CNN Dropout=0.1: Epochs=20, Train accuracy=0.99852, Validation accuracy=0.99160
 CNN Dropout=0.2: Epochs=20, Train accuracy=0.99865, Validation accuracy=0.99290
 CNN Dropout=0.3: Epochs=20, Train accuracy=0.99782, Validation accuracy=0.99240
 CNN Dropout=0.4: Epochs=20, Train accuracy=0.99820, Validation accuracy=0.99130
 CNN Dropout=0.5: Epochs=20, Train accuracy=0.99792, Validation accuracy=0.99090
 CNN Dropout=0.6: Epochs=20, Train accuracy=0.99817, Validation accuracy=0.99080
 CNN Dropout=0.7: Epochs=20, Train accuracy=0.99783, Validation accuracy=0.99270

7-Layer ConvNet + Adam + Relu

Batch Normalization with Dropout rate 0.2

```
In [19]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20

#ConvNet
mdl_5 = Sequential()
mdl_5.add(Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=(input_shape)))
mdl_5.add(Conv2D(32, kernel_size=(3,3), activation='relu'))
mdl_5.add(Conv2D(32, kernel_size=(3,3), activation='relu'))
mdl_5.add(Dropout(rate=0.2))

mdl_5.add(Conv2D(64, kernel_size=(3,3), activation='relu'))
mdl_5.add(Conv2D(64, kernel_size=(3,3), activation='relu'))

mdl_5.add(Dropout(rate=0.2))
mdl_5.add(BatchNormalization())

mdl_5.add(Conv2D(64, kernel_size=(5,5), activation='relu'))
mdl_5.add(Conv2D(128, kernel_size=(2,2), activation='relu'))

mdl_5.add(BatchNormalization())
mdl_5.add(Flatten())
mdl_5.add(Dense(128, activation='relu'))
mdl_5.add(Dense(num_clas, activation='softmax'))

mdl_5.summary()

mdl_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history_5 = mdl_5.fit(X_train, y_train, batch_size=batch_size, epochs=no_epoch, verbose=1, validation_data=(x_test, y_test))
```


Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_22 (Conv2D)	(None, 26, 26, 32)	320
conv2d_23 (Conv2D)	(None, 24, 24, 32)	9248
conv2d_24 (Conv2D)	(None, 22, 22, 32)	9248
dropout_7 (Dropout)	(None, 22, 22, 32)	0
conv2d_25 (Conv2D)	(None, 20, 20, 64)	18496
conv2d_26 (Conv2D)	(None, 18, 18, 64)	36928
dropout_8 (Dropout)	(None, 18, 18, 64)	0
batch_normalization_7 (Batch Normalization)	(None, 18, 18, 64)	256
conv2d_27 (Conv2D)	(None, 14, 14, 64)	102464
conv2d_28 (Conv2D)	(None, 13, 13, 128)	32896
batch_normalization_8 (Batch Normalization)	(None, 13, 13, 128)	512
flatten_4 (Flatten)	(None, 21632)	0
dense_7 (Dense)	(None, 128)	2769024
dense_8 (Dense)	(None, 10)	1290
Total params: 2,980,682		
Trainable params: 2,980,298		
Non-trainable params: 384		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 30s 498us/step - loss: 0.1905 - acc: 0.9527 - val_loss: 0.1044 - val_acc: 0.9759

Epoch 2/20

60000/60000 [=====] - 28s 473us/step - loss: 0.0632 - acc: 0.9831 - val_loss: 0.0770 - val_acc: 0.9785

Epoch 3/20

60000/60000 [=====] - 28s 470us/step - loss: 0.0423 - acc: 0.9881 - val_loss: 0.0634 - val_acc: 0.9839

Epoch 4/20

60000/60000 [=====] - 28s 473us/step - loss: 0.0304 - acc: 0.9916 - val_loss: 0.0506 - val_acc: 0.9887

Epoch 5/20

60000/60000 [=====] - 28s 472us/step - loss: 0.0242 - acc: 0.9928 - val_loss: 0.0525 - val_acc: 0.9858

Epoch 6/20

60000/60000 [=====] - 28s 470us/step - loss: 0.0269 - acc: 0.9927 - val_loss: 0.0484 - val_acc: 0.9876

Epoch 7/20

60000/60000 [=====] - 28s 471us/step - loss: 0.022


5 - acc: 0.9935 - val_loss: 0.0672 - val_acc: 0.9865
Epoch 8/20
60000/60000 [=====] - 28s 472us/step - loss: 0.022
4 - acc: 0.9935 - val_loss: 0.0628 - val_acc: 0.9876
Epoch 9/20
60000/60000 [=====] - 28s 472us/step - loss: 0.016
6 - acc: 0.9952 - val_loss: 0.0540 - val_acc: 0.9886
Epoch 10/20
60000/60000 [=====] - 28s 471us/step - loss: 0.017
6 - acc: 0.9949 - val_loss: 0.0542 - val_acc: 0.9900
Epoch 11/20
60000/60000 [=====] - 28s 472us/step - loss: 0.015
8 - acc: 0.9957 - val_loss: 0.0635 - val_acc: 0.9866
Epoch 12/20
60000/60000 [=====] - 28s 472us/step - loss: 0.014
5 - acc: 0.9959 - val_loss: 0.0572 - val_acc: 0.9879
Epoch 13/20
60000/60000 [=====] - 29s 476us/step - loss: 0.013
7 - acc: 0.9962 - val_loss: 0.0562 - val_acc: 0.9898
Epoch 14/20
60000/60000 [=====] - 29s 475us/step - loss: 0.010
5 - acc: 0.9973 - val_loss: 0.0767 - val_acc: 0.9854
Epoch 15/20
60000/60000 [=====] - 29s 477us/step - loss: 0.012
2 - acc: 0.9969 - val_loss: 0.0505 - val_acc: 0.9898
Epoch 16/20
60000/60000 [=====] - 29s 480us/step - loss: 0.010
7 - acc: 0.9971 - val_loss: 0.0530 - val_acc: 0.9904
Epoch 17/20
60000/60000 [=====] - 28s 472us/step - loss: 0.009
0 - acc: 0.9975 - val_loss: 0.0560 - val_acc: 0.9906
Epoch 18/20
60000/60000 [=====] - 28s 472us/step - loss: 0.013
2 - acc: 0.9965 - val_loss: 0.0778 - val_acc: 0.9850
Epoch 19/20
60000/60000 [=====] - 28s 472us/step - loss: 0.009
9 - acc: 0.9975 - val_loss: 0.0474 - val_acc: 0.9917
Epoch 20/20
60000/60000 [=====] - 28s 471us/step - loss: 0.006
8 - acc: 0.9981 - val_loss: 0.0517 - val_acc: 0.9919

```
In [20]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_5.history['loss']
test_loss = history_5.history['val_loss']

# Create count of the number of epochs
score = mdl_5.evaluate(x_test, y_test, verbose=0)
plt_mdl_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.05174927458126408
Test accuracy: 0.9919

Conclusion

```
In [3]: import tabulate
res_tab = [['Layer','Batch \nNormalization','Dropout','Dropout \nRate','Tes
t \nAccuracy'],
           [3,'No','No',0.0,0.9938],
           [3,'No','Yes',0.0,0.9928],
           [5,'Yes','No',0.0,0.9934],
           [7,'Yes','Yes',0.0,0.9810],
           [7,'Yes','Yes',0.2,0.9919]
          ]
print(tabulate.tabulate(res_tab,tablefmt='fancy_grid'))
```

Layer	Batch Normalization	Dropout	Dropout Rate	Test Accuracy
3	No	No	0.0	0.9938
3	No	Yes	0.0	0.9928
5	Yes	No	0.0	0.9934
7	Yes	Yes	0.0	0.981
7	Yes	Yes	0.2	0.9919