

## Keras -- MLPs on MNIST

```
In [1]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensor
flow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
from keras.initializers import RandomNormal
from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
```

Using TensorFlow backend.

```
In [0]: import seaborn as sns
```

```
In [0]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [3]: # the data, shuffled and split between train and test sets
(X_train, y_train), (x_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 0s 0us/step

```
In [4]: print("Number of training examples :", X_train.shape[0], X_train.shape[1],
X_train.shape[2])
```

Number of training examples : 60000 28 28

```
In [5]: print("Number of training examples :", X_train.shape[0], "and each image is
of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", x_test.shape[0], "and each image is
of shape (%d, %d)"%(x_test.shape[1], x_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 *
784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[
2])
x_test = x_test.reshape(x_test.shape[0], x_test.shape[1]*x_test.shape[2])
```

```
In [8]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is
of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", x_test.shape[0], "and each image is
of shape (%d)"%(x_test.shape[1]))
```

```
Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)
```

```
In [9]: # An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0 249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
0]
```

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize
the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
x_test = x_test/255
```

```
In [0]: # example data point after normlizing  
print(X_train[0])
```

file:///C:/Users/Admin/AppData/Local/Temp/ariyurjana@gmail.com\_12.html

0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451



```

In [0]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to
the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initiali
zer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) +
bias) where
# activation is the element-wise activation function passed as the activati
on argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias i
s True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X +
b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through th
e activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

```



```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

```
In [0]: # start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

```
In [0]: # Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer, https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted our labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples  
Epoch 1/20  
60000/60000 [=====] - 3s 49us/step - loss: 1.2935  
- acc: 0.6829 - val\_loss: 0.8171 - val\_acc: 0.8312  
Epoch 2/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.7213  
- acc: 0.8385 - val\_loss: 0.6099 - val\_acc: 0.8623  
Epoch 3/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.5904  
- acc: 0.8584 - val\_loss: 0.5275 - val\_acc: 0.8741  
Epoch 4/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.5278  
- acc: 0.8682 - val\_loss: 0.4815 - val\_acc: 0.8814  
Epoch 5/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.4897  
- acc: 0.8747 - val\_loss: 0.4515 - val\_acc: 0.8870  
Epoch 6/20  
34432/60000 [=====>.....] - ETA: 0s - loss: 0.4697 - acc: 0.8773  
60000/60000 [=====] - 2s 35us/step - loss: 0.4635 - acc: 0.8799 - val\_loss: 0.4303 - val\_acc: 0.8898  
Epoch 7/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.4442  
- acc: 0.8837 - val\_loss: 0.4138 - val\_acc: 0.8917  
Epoch 8/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.4290  
- acc: 0.8866 - val\_loss: 0.4010 - val\_acc: 0.8952  
Epoch 9/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.4169  
- acc: 0.8894 - val\_loss: 0.3909 - val\_acc: 0.8971  
Epoch 10/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.4067  
- acc: 0.8911 - val\_loss: 0.3818 - val\_acc: 0.8994  
Epoch 11/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.3982  
- acc: 0.8926 - val\_loss: 0.3743 - val\_acc: 0.9006  
Epoch 12/20  
1792/60000 [.....] - ETA: 1s - loss: 0.4077 - acc: 0.8895  
60000/60000 [=====] - 2s 35us/step - loss: 0.3908 - acc: 0.8948 - val\_loss: 0.3681 - val\_acc: 0.9020  
Epoch 13/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.3844  
- acc: 0.8960 - val\_loss: 0.3624 - val\_acc: 0.9039  
Epoch 14/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.3786  
- acc: 0.8972 - val\_loss: 0.3575 - val\_acc: 0.9046  
Epoch 15/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.3735  
- acc: 0.8981 - val\_loss: 0.3528 - val\_acc: 0.9058  
Epoch 16/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.3689  
- acc: 0.8993 - val\_loss: 0.3490 - val\_acc: 0.9062  
Epoch 17/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.3648  
- acc: 0.9004 - val\_loss: 0.3455 - val\_acc: 0.9066  
Epoch 18/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.3610  
- acc: 0.9016 - val\_loss: 0.3419 - val\_acc: 0.9077

```
Epoch 19/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3575
- acc: 0.9024 - val_loss: 0.3389 - val_acc: 0.9088
Epoch 20/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3544
- acc: 0.9032 - val_loss: 0.3362 - val_acc: 0.9093
```

```
In [0]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=
nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validati
on_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

Test score: 0.3362289469957352
Test accuracy: 0.9093
```

## MLP + Sigmoid activation + SGDOptimizer

```
In [0]: # Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
In [0]: model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20  
60000/60000 [=====] - 3s 42us/step - loss: 2.2708  
- acc: 0.2169 - val\_loss: 2.2197 - val\_acc: 0.2768

Epoch 2/20  
60000/60000 [=====] - 2s 42us/step - loss: 2.1739  
- acc: 0.4237 - val\_loss: 2.1143 - val\_acc: 0.4877

Epoch 3/20  
60000/60000 [=====] - 2s 42us/step - loss: 2.0506  
- acc: 0.5485 - val\_loss: 1.9659 - val\_acc: 0.5524

Epoch 4/20  
60000/60000 [=====] - 3s 42us/step - loss: 1.8796  
- acc: 0.6135 - val\_loss: 1.7673 - val\_acc: 0.6481

Epoch 5/20  
60000/60000 [=====] - 3s 42us/step - loss: 1.6674  
- acc: 0.6659 - val\_loss: 1.5414 - val\_acc: 0.7205

Epoch 6/20  
5376/60000 [=>.....] - ETA: 2s - loss: 1.5430 - acc: 0.698160000/60000 [=====] - 2s 41us/step - loss: 1.4449 - acc: 0.7125 - val\_loss: 1.3233 - val\_acc: 0.7513

Epoch 7/20  
60000/60000 [=====] - 2s 41us/step - loss: 1.2442  
- acc: 0.7466 - val\_loss: 1.1406 - val\_acc: 0.7599

Epoch 8/20  
60000/60000 [=====] - 2s 41us/step - loss: 1.0815  
- acc: 0.7722 - val\_loss: 0.9974 - val\_acc: 0.7968

Epoch 9/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.9544  
- acc: 0.7940 - val\_loss: 0.8867 - val\_acc: 0.8060

Epoch 10/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.8556  
- acc: 0.8082 - val\_loss: 0.7984 - val\_acc: 0.8227

Epoch 11/20  
31232/60000 [=====>.....] - ETA: 1s - loss: 0.7920 - acc: 0.820760000/60000 [=====] - 2s 42us/step - loss: 0.7776 - acc: 0.8225 - val\_loss: 0.7292 - val\_acc: 0.8335

Epoch 12/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.7154  
- acc: 0.8333 - val\_loss: 0.6741 - val\_acc: 0.8422

Epoch 13/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.6650  
- acc: 0.8412 - val\_loss: 0.6282 - val\_acc: 0.8500

Epoch 14/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.6237  
- acc: 0.8485 - val\_loss: 0.5906 - val\_acc: 0.8585

Epoch 15/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.5893  
- acc: 0.8546 - val\_loss: 0.5591 - val\_acc: 0.8616

Epoch 16/20  
34432/60000 [=====>.....] - ETA: 0s - loss: 0.5691 - acc: 0.857860000/60000 [=====] - 2s 41us/step - loss: 0.5606 - acc: 0.8599 - val\_loss: 0.5329 - val\_acc: 0.8672

Epoch 17/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.5361  
- acc: 0.8641 - val\_loss: 0.5095 - val\_acc: 0.8703

Epoch 18/20  
60000/60000 [=====] - 2s 42us/step - loss: 0.5151

```
- acc: 0.8676 - val_loss: 0.4904 - val_acc: 0.8736
Epoch 19/20
60000/60000 [=====] - 2s 41us/step - loss: 0.4969
- acc: 0.8710 - val_loss: 0.4732 - val_acc: 0.8782
Epoch 20/20
60000/60000 [=====] - 2s 42us/step - loss: 0.4810
- acc: 0.8739 - val_loss: 0.4583 - val_acc: 0.8816
```

```
In [0]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=
nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validati
on_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

Test score: 0.4582893396139145
Test accuracy: 0.8816
```



```
In [0]: w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)
```

## MLP + Sigmoid activation + ADAM

```
In [0]: model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 55us/step - loss: 0.5308  
- acc: 0.8636 - val\_loss: 0.2550 - val\_acc: 0.9259

Epoch 2/20

53120/60000 [=====>....] - ETA: 0s - loss: 0.2244 - acc: 0.9340  
60000/60000 [=====] - 3s 51us/step - loss: 0.2205 - acc: 0.9351 - val\_loss: 0.1946 - val\_acc: 0.9417

Epoch 3/20

60000/60000 [=====] - 3s 51us/step - loss: 0.1650  
- acc: 0.9512 - val\_loss: 0.1421 - val\_acc: 0.9570

Epoch 4/20

60000/60000 [=====] - 3s 51us/step - loss: 0.1279  
- acc: 0.9614 - val\_loss: 0.1238 - val\_acc: 0.9645

Epoch 5/20

60000/60000 [=====] - 3s 51us/step - loss: 0.1010  
- acc: 0.9704 - val\_loss: 0.1029 - val\_acc: 0.9693

Epoch 6/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0801  
- acc: 0.9763 - val\_loss: 0.0877 - val\_acc: 0.9725

Epoch 7/20

4480/60000 [=>.....] - ETA: 2s - loss: 0.0632 - acc: 0.9795  
60000/60000 [=====] - 3s 51us/step - loss: 0.0645 - acc: 0.9809 - val\_loss: 0.0831 - val\_acc: 0.9751

Epoch 8/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0519  
- acc: 0.9842 - val\_loss: 0.0724 - val\_acc: 0.9780

Epoch 9/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0433  
- acc: 0.9872 - val\_loss: 0.0714 - val\_acc: 0.9786

Epoch 10/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0347  
- acc: 0.9898 - val\_loss: 0.0695 - val\_acc: 0.9776

Epoch 11/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0268  
- acc: 0.9930 - val\_loss: 0.0659 - val\_acc: 0.9796

Epoch 12/20

60000/60000 [=====] - 3s 52us/step - loss: 0.0219  
- acc: 0.9944 - val\_loss: 0.0642 - val\_acc: 0.9809

Epoch 13/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0180  
- acc: 0.9953 - val\_loss: 0.0677 - val\_acc: 0.9794

Epoch 14/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0133

```
- acc: 0.9970 - val_loss: 0.0647 - val_acc: 0.9803
Epoch 15/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0114
- acc: 0.9975 - val_loss: 0.0628 - val_acc: 0.9812
Epoch 16/20
57344/60000 [=====>..] - ETA: 0s - loss: 0.0085 - acc: 0.9982
60000/60000 [=====] - 3s 50us/step - loss: 0.0085 - acc: 0.9982 - val_loss: 0.0666 - val_acc: 0.9806
Epoch 17/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0070
- acc: 0.9986 - val_loss: 0.0643 - val_acc: 0.9822
Epoch 18/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0061
- acc: 0.9986 - val_loss: 0.0656 - val_acc: 0.9818
Epoch 19/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0055
- acc: 0.9988 - val_loss: 0.0811 - val_acc: 0.9774
Epoch 20/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0038
- acc: 0.9992 - val_loss: 0.0723 - val_acc: 0.9818
```

```
In [0]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=
nb_epoch, verbose=1, validation_data=(X_test, Y_test))

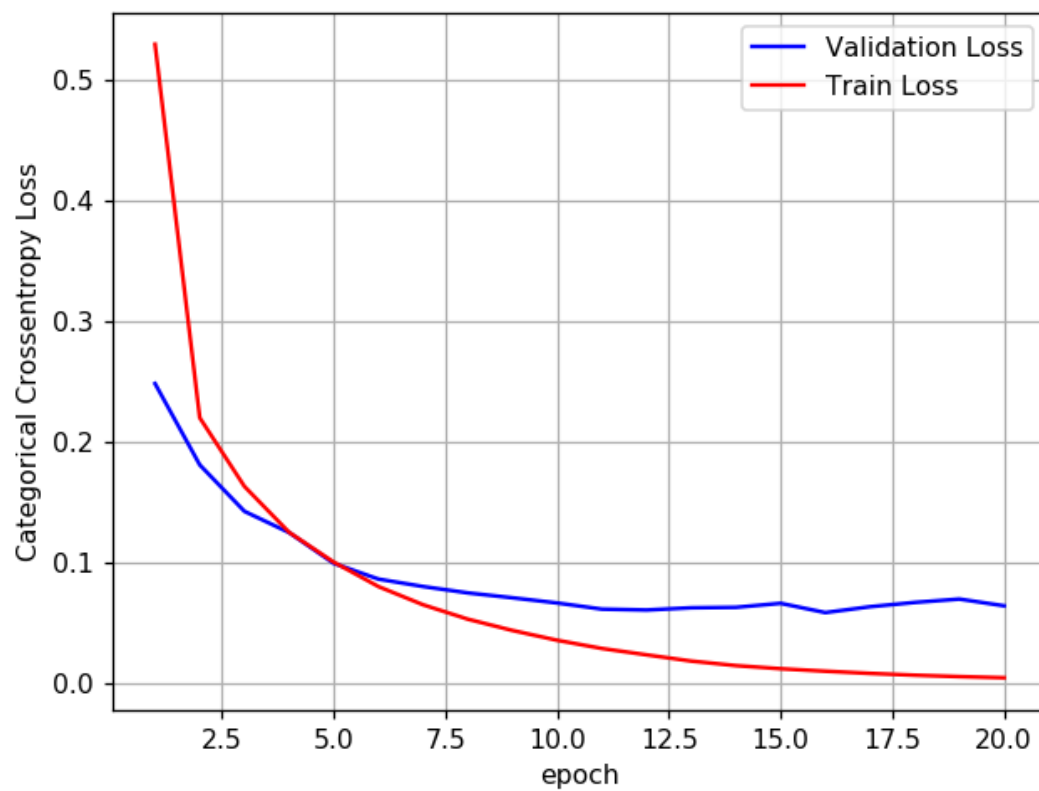
# we will get val_loss and val_acc only when you pass the paramter validati
on_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06385514608082886

Test accuracy: 0.9824



```
In [0]: w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)
```

## MLP + ReLU +SGD

```
In [0]: # Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		



```
In [0]: model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20  
60000/60000 [=====] - 4s 67us/step - loss: 0.7579  
- acc: 0.7812 - val\_loss: 0.3951 - val\_acc: 0.8921

Epoch 2/20  
60000/60000 [=====] - 4s 64us/step - loss: 0.3535  
- acc: 0.8998 - val\_loss: 0.3040 - val\_acc: 0.9153

Epoch 3/20  
60000/60000 [=====] - 4s 64us/step - loss: 0.2900  
- acc: 0.9172 - val\_loss: 0.2648 - val\_acc: 0.9253

Epoch 4/20  
60000/60000 [=====] - 4s 60us/step - loss: 0.2558  
- acc: 0.9269 - val\_loss: 0.2393 - val\_acc: 0.9316

Epoch 5/20  
60000/60000 [=====] - 4s 58us/step - loss: 0.2324  
- acc: 0.9340 - val\_loss: 0.2210 - val\_acc: 0.9371

Epoch 6/20  
60000/60000 [=====] - 4s 64us/step - loss: 0.2144  
- acc: 0.9391 - val\_loss: 0.2072 - val\_acc: 0.9400

Epoch 7/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.1995  
- acc: 0.9443 - val\_loss: 0.1957 - val\_acc: 0.9444

Epoch 8/20  
60000/60000 [=====] - 4s 61us/step - loss: 0.1872  
- acc: 0.9476 - val\_loss: 0.1848 - val\_acc: 0.9456

Epoch 9/20  
60000/60000 [=====] - 3s 57us/step - loss: 0.1763  
- acc: 0.9507 - val\_loss: 0.1771 - val\_acc: 0.9488

Epoch 10/20  
60000/60000 [=====] - 4s 60us/step - loss: 0.1668  
- acc: 0.9539 - val\_loss: 0.1682 - val\_acc: 0.9506

Epoch 11/20  
60000/60000 [=====] - 4s 71us/step - loss: 0.1585  
- acc: 0.9560 - val\_loss: 0.1623 - val\_acc: 0.9518

Epoch 12/20  
60000/60000 [=====] - 7s 113us/step - loss: 0.1511  
- acc: 0.9577 - val\_loss: 0.1560 - val\_acc: 0.9543

Epoch 13/20  
60000/60000 [=====] - 7s 115us/step - loss: 0.1443  
- acc: 0.9596 - val\_loss: 0.1517 - val\_acc: 0.9557

Epoch 14/20  
60000/60000 [=====] - 7s 111us/step - loss: 0.1379  
- acc: 0.9615 - val\_loss: 0.1474 - val\_acc: 0.9572

Epoch 15/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.1323  
- acc: 0.9628 - val\_loss: 0.1429 - val\_acc: 0.9580

Epoch 16/20  
60000/60000 [=====] - 4s 69us/step - loss: 0.1270  
- acc: 0.9645 - val\_loss: 0.1371 - val\_acc: 0.9598

Epoch 17/20  
60000/60000 [=====] - 7s 110us/step - loss: 0.1221  
- acc: 0.9661 - val\_loss: 0.1351 - val\_acc: 0.9602

Epoch 18/20  
60000/60000 [=====] - 4s 62us/step - loss: 0.1177  
- acc: 0.9671 - val\_loss: 0.1309 - val\_acc: 0.9618

Epoch 19/20  
60000/60000 [=====] - 4s 60us/step - loss: 0.1136

- acc: 0.9685 - val\_loss: 0.1263 - val\_acc: 0.9631

Epoch 20/20

60000/60000 [=====] - 5s 79us/step - loss: 0.1094

- acc: 0.9694 - val\_loss: 0.1241 - val\_acc: 0.9631

```
In [0]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=
nb_epoch, verbose=1, validation_data=(X_test, Y_test))

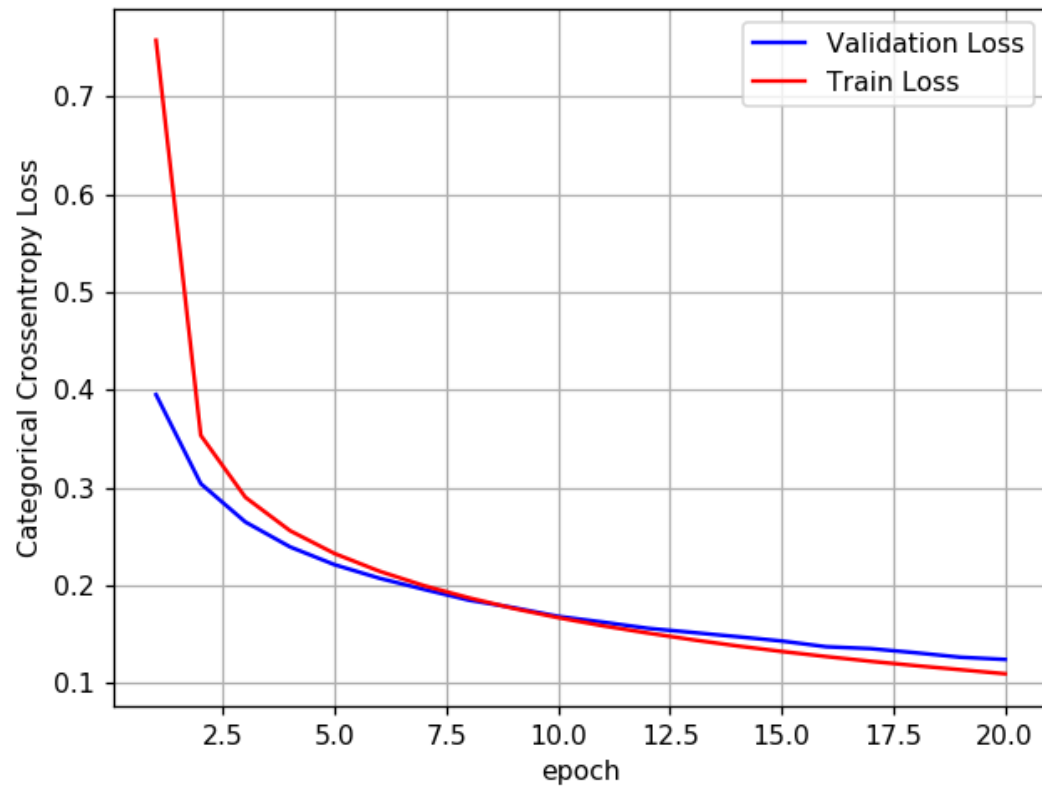
# we will get val_loss and val_acc only when you pass the paramter validati
on_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12405014228336513

Test accuracy: 0.9631



```

In [0]: w_after = model_relu.get_weights()

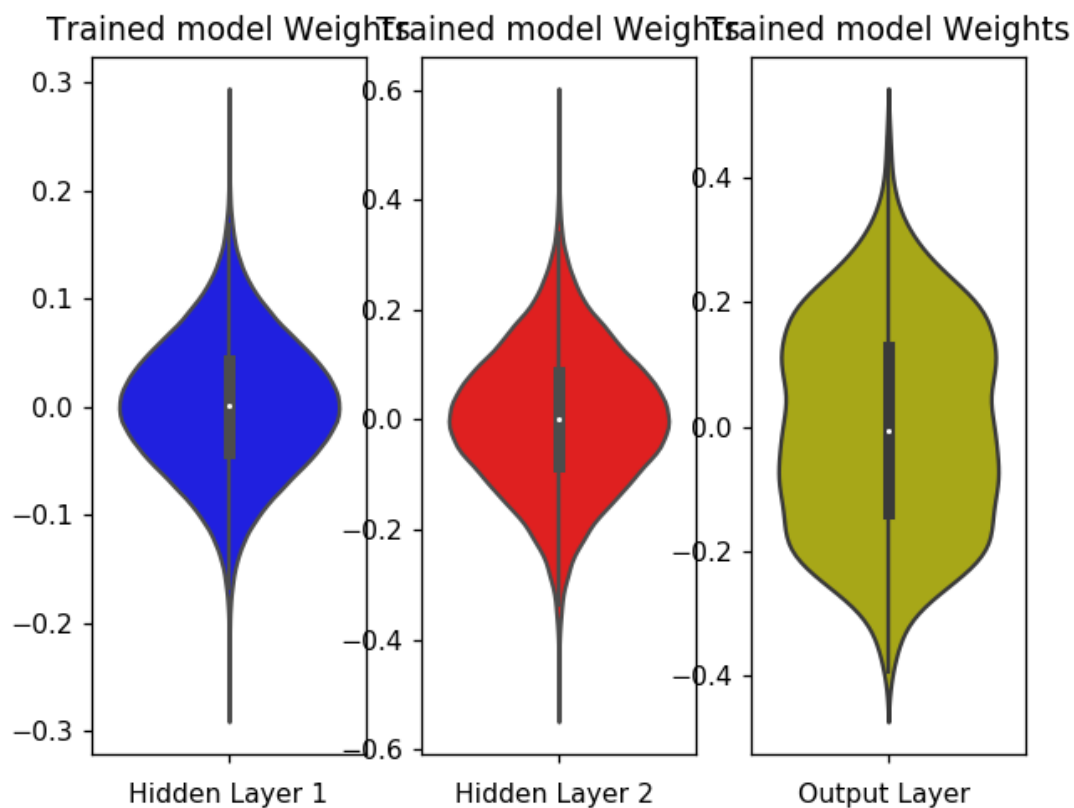
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## MLP + ReLU + ADAM

```
In [0]: model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```



Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290

=====  
 Total params: 468,874  
 Trainable params: 468,874  
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 121us/step - loss: 0.2341  
 - acc: 0.9295 - val\_loss: 0.1165 - val\_acc: 0.9652

Epoch 2/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0878  
 - acc: 0.9729 - val\_loss: 0.0883 - val\_acc: 0.9720

Epoch 3/20

60000/60000 [=====] - 5s 75us/step - loss: 0.0544  
 - acc: 0.9825 - val\_loss: 0.0860 - val\_acc: 0.9729

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0354  
 - acc: 0.9885 - val\_loss: 0.0699 - val\_acc: 0.9797

Epoch 5/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0266  
 - acc: 0.9914 - val\_loss: 0.0720 - val\_acc: 0.9788

Epoch 6/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0200  
 - acc: 0.9941 - val\_loss: 0.0696 - val\_acc: 0.9803

Epoch 7/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0155  
 - acc: 0.9951 - val\_loss: 0.0640 - val\_acc: 0.9829

Epoch 8/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0140  
 - acc: 0.9952 - val\_loss: 0.0848 - val\_acc: 0.9792

Epoch 9/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0143  
 - acc: 0.9952 - val\_loss: 0.0837 - val\_acc: 0.9796

Epoch 10/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0128  
 - acc: 0.9958 - val\_loss: 0.0946 - val\_acc: 0.9782

Epoch 11/20

60000/60000 [=====] - 7s 125us/step - loss: 0.0081  
 - acc: 0.9974 - val\_loss: 0.0682 - val\_acc: 0.9826

Epoch 12/20

60000/60000 [=====] - 8s 129us/step - loss: 0.0121  
 - acc: 0.9959 - val\_loss: 0.0793 - val\_acc: 0.9816

Epoch 13/20

60000/60000 [=====] - 8s 133us/step - loss: 0.0107  
 - acc: 0.9963 - val\_loss: 0.0746 - val\_acc: 0.9820

Epoch 14/20

60000/60000 [=====] - 8s 129us/step - loss: 0.0113  
 - acc: 0.9960 - val\_loss: 0.0813 - val\_acc: 0.9816

Epoch 15/20  
60000/60000 [=====] - 5s 77us/step - loss: 0.0058  
- acc: 0.9982 - val\_loss: 0.0770 - val\_acc: 0.9842

Epoch 16/20  
60000/60000 [=====] - 4s 65us/step - loss: 0.0040  
- acc: 0.9987 - val\_loss: 0.0930 - val\_acc: 0.9808

Epoch 17/20  
60000/60000 [=====] - 4s 68us/step - loss: 0.0119  
- acc: 0.9959 - val\_loss: 0.0813 - val\_acc: 0.9819

Epoch 18/20  
60000/60000 [=====] - 4s 73us/step - loss: 0.0105  
- acc: 0.9966 - val\_loss: 0.1000 - val\_acc: 0.9803

Epoch 19/20  
60000/60000 [=====] - 4s 69us/step - loss: 0.0064  
- acc: 0.9981 - val\_loss: 0.0852 - val\_acc: 0.9831

Epoch 20/20  
60000/60000 [=====] - 4s 72us/step - loss: 0.0056  
- acc: 0.9982 - val\_loss: 0.1029 - val\_acc: 0.9805

```
In [0]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=
nb_epoch, verbose=1, validation_data=(X_test, Y_test))

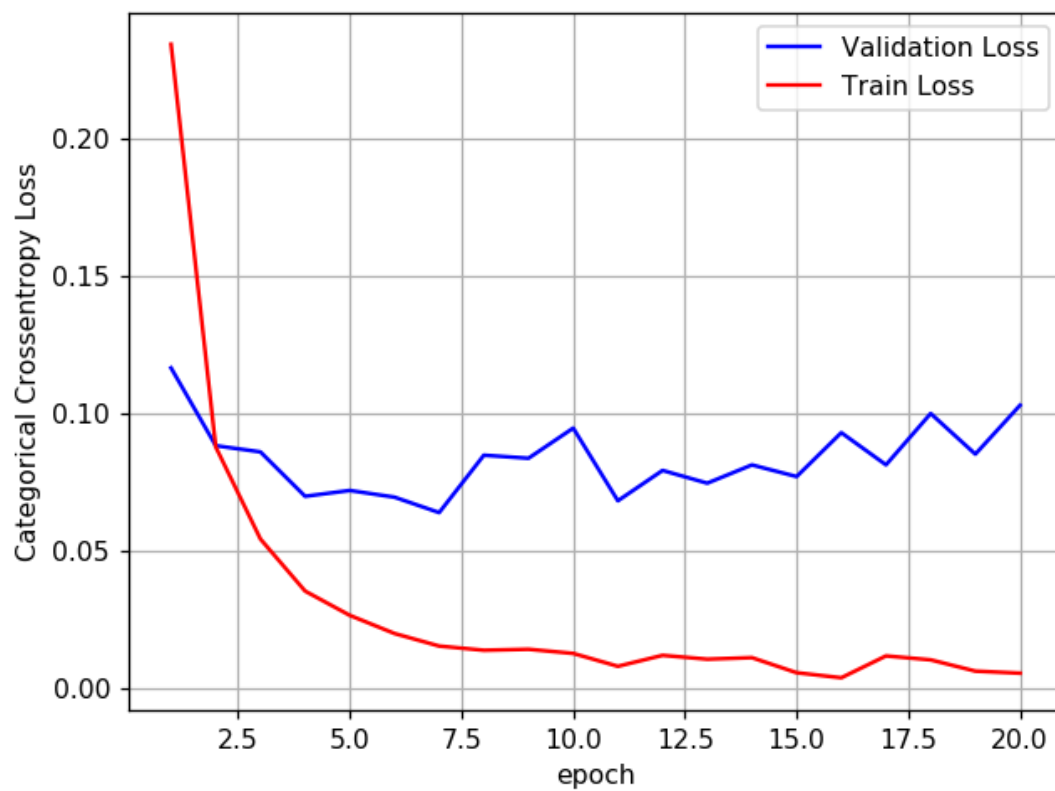
# we will get val_loss and val_acc only when you pass the paramter validati
on_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10294274219236926

Test accuracy: 0.9805



```

In [0]: w_after = model_relu.get_weights()

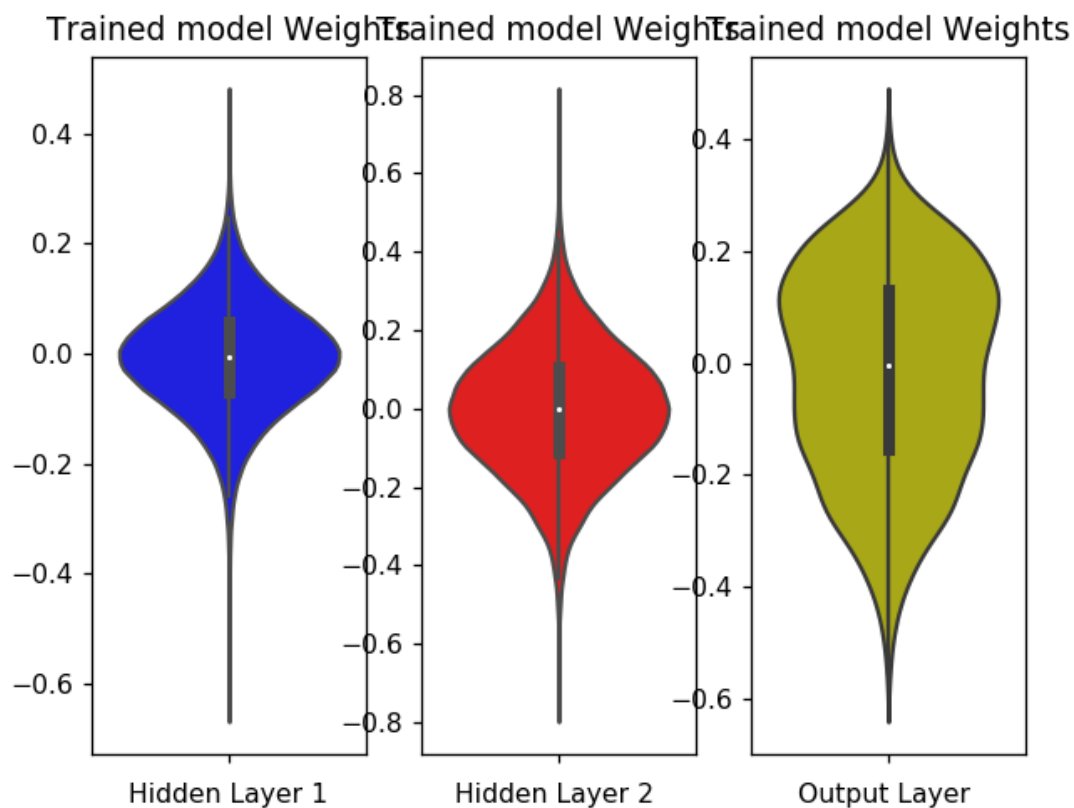
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

```
In [0]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_14 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		
=====	=====	=====

```
In [0]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=n_b_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20  
60000/60000 [=====] - 8s 138us/step - loss: 0.3036  
- acc: 0.9104 - val\_loss: 0.2116 - val\_acc: 0.9376

Epoch 2/20  
60000/60000 [=====] - 10s 170us/step - loss: 0.1747  
- acc: 0.9483 - val\_loss: 0.1670 - val\_acc: 0.9505

Epoch 3/20  
60000/60000 [=====] - 13s 220us/step - loss: 0.1367  
- acc: 0.9599 - val\_loss: 0.1451 - val\_acc: 0.9567

Epoch 4/20  
60000/60000 [=====] - 9s 156us/step - loss: 0.1134  
- acc: 0.9666 - val\_loss: 0.1335 - val\_acc: 0.9603

Epoch 5/20  
60000/60000 [=====] - 13s 211us/step - loss: 0.0949  
- acc: 0.9703 - val\_loss: 0.1325 - val\_acc: 0.9589

Epoch 6/20  
60000/60000 [=====] - 7s 119us/step - loss: 0.0802  
- acc: 0.9758 - val\_loss: 0.1139 - val\_acc: 0.9652

Epoch 7/20  
60000/60000 [=====] - 8s 127us/step - loss: 0.0682  
- acc: 0.9787 - val\_loss: 0.1136 - val\_acc: 0.9666

Epoch 8/20  
60000/60000 [=====] - 7s 124us/step - loss: 0.0608  
- acc: 0.9815 - val\_loss: 0.1114 - val\_acc: 0.9666

Epoch 9/20  
60000/60000 [=====] - 8s 129us/step - loss: 0.0532  
- acc: 0.9837 - val\_loss: 0.1167 - val\_acc: 0.9666

Epoch 10/20  
60000/60000 [=====] - 7s 123us/step - loss: 0.0455  
- acc: 0.9856 - val\_loss: 0.0962 - val\_acc: 0.9718

Epoch 11/20  
60000/60000 [=====] - 7s 112us/step - loss: 0.0376  
- acc: 0.9880 - val\_loss: 0.1102 - val\_acc: 0.9673

Epoch 12/20  
60000/60000 [=====] - 7s 124us/step - loss: 0.0350  
- acc: 0.9889 - val\_loss: 0.1033 - val\_acc: 0.9710

Epoch 13/20  
60000/60000 [=====] - 7s 124us/step - loss: 0.0308  
- acc: 0.9903 - val\_loss: 0.1020 - val\_acc: 0.9712

Epoch 14/20  
60000/60000 [=====] - 7s 123us/step - loss: 0.0271  
- acc: 0.9913 - val\_loss: 0.1038 - val\_acc: 0.9727

Epoch 15/20  
60000/60000 [=====] - 7s 122us/step - loss: 0.0231  
- acc: 0.9926 - val\_loss: 0.1019 - val\_acc: 0.9717

Epoch 16/20  
60000/60000 [=====] - 8s 127us/step - loss: 0.0220  
- acc: 0.9928 - val\_loss: 0.1110 - val\_acc: 0.9703

Epoch 17/20  
60000/60000 [=====] - 7s 114us/step - loss: 0.0229  
- acc: 0.9928 - val\_loss: 0.1067 - val\_acc: 0.9739

Epoch 18/20  
60000/60000 [=====] - 8s 128us/step - loss: 0.0203  
- acc: 0.9935 - val\_loss: 0.0982 - val\_acc: 0.9738

Epoch 19/20  
60000/60000 [=====] - 7s 125us/step - loss: 0.0171



- acc: 0.9944 - val\_loss: 0.1056 - val\_acc: 0.9706

Epoch 20/20

60000/60000 [=====] - 11s 182us/step - loss: 0.014

6 - acc: 0.9952 - val\_loss: 0.1046 - val\_acc: 0.9732

```
In [0]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=
nb_epoch, verbose=1, validation_data=(X_test, Y_test))

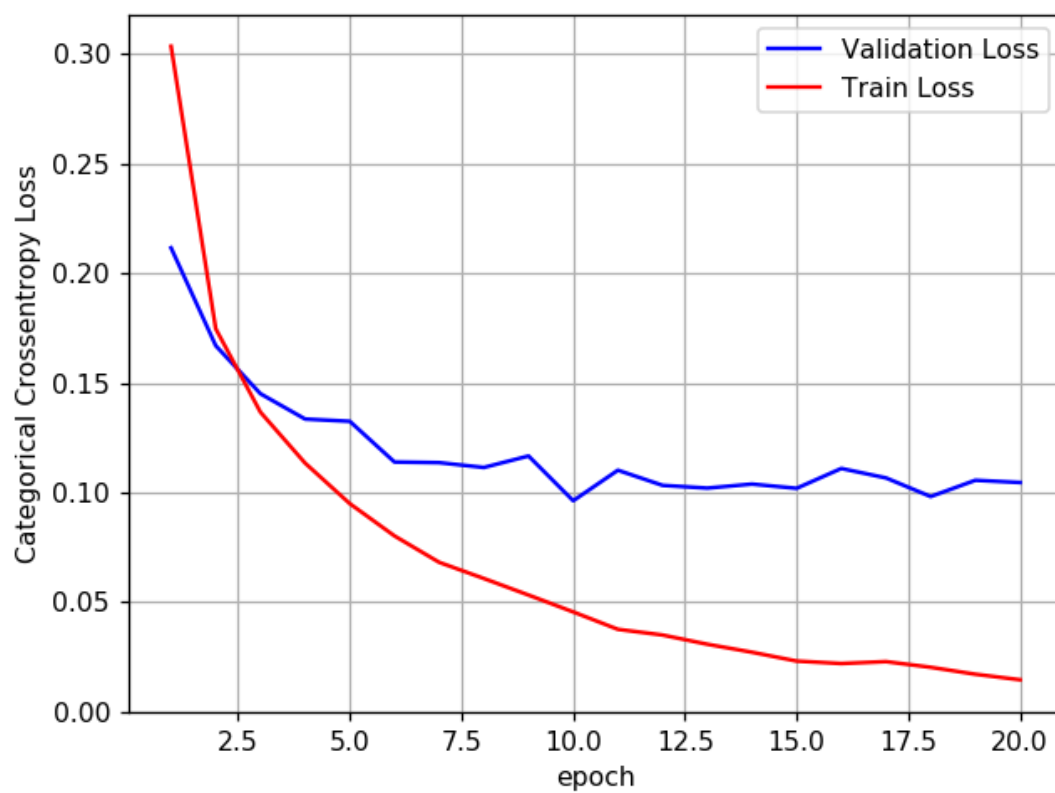
# we will get val_loss and val_acc only when you pass the paramter validati
on_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10456635547156475

Test accuracy: 0.9732



```

In [0]: w_after = model_batch.get_weights()

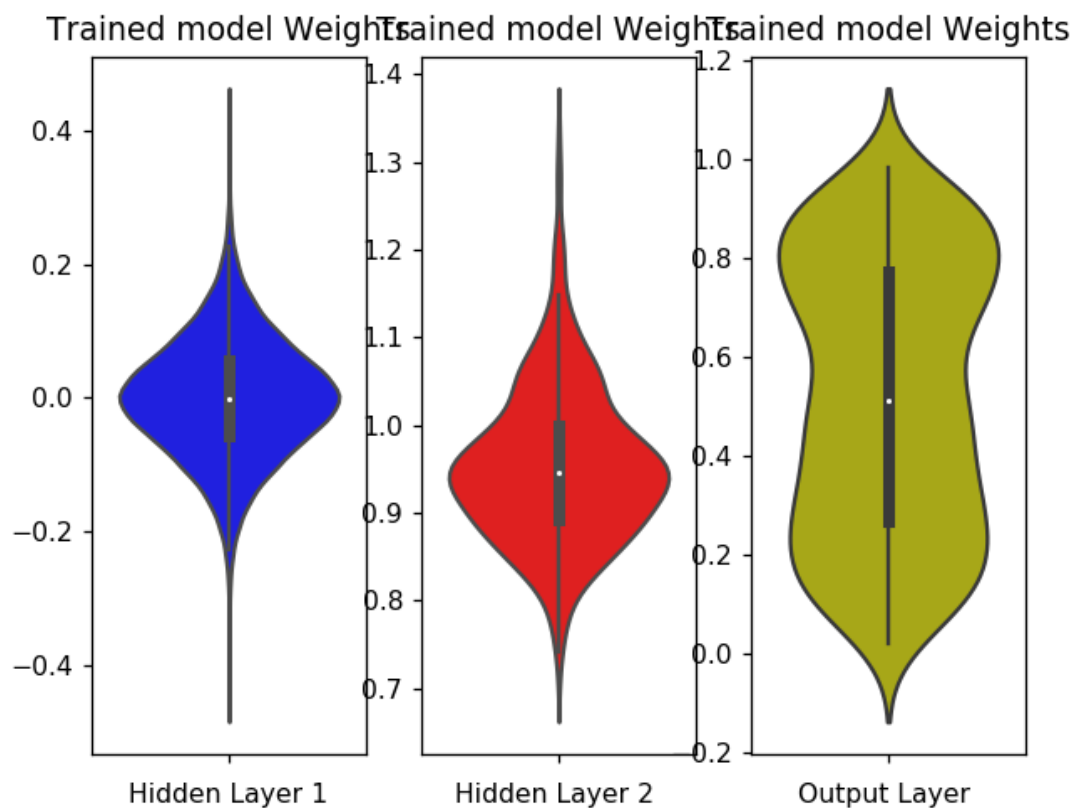
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## 5. MLP + Dropout + AdamOptimizer

```
In [0]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_17 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

```
In [0]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20  
60000/60000 [=====] - 14s 227us/step - loss: 0.661  
2 - acc: 0.7951 - val\_loss: 0.2860 - val\_acc: 0.9166

Epoch 2/20  
60000/60000 [=====] - 8s 136us/step - loss: 0.4250  
- acc: 0.8710 - val\_loss: 0.2545 - val\_acc: 0.9252

Epoch 3/20  
60000/60000 [=====] - 12s 198us/step - loss: 0.384  
1 - acc: 0.8846 - val\_loss: 0.2391 - val\_acc: 0.9298

Epoch 4/20  
60000/60000 [=====] - 8s 138us/step - loss: 0.3551  
- acc: 0.8927 - val\_loss: 0.2279 - val\_acc: 0.9325

Epoch 5/20  
60000/60000 [=====] - 7s 123us/step - loss: 0.3355  
- acc: 0.8986 - val\_loss: 0.2127 - val\_acc: 0.9356

Epoch 6/20  
60000/60000 [=====] - 8s 136us/step - loss: 0.3234  
- acc: 0.9031 - val\_loss: 0.2029 - val\_acc: 0.9387: 1s - loss:

Epoch 7/20  
60000/60000 [=====] - 8s 131us/step - loss: 0.3068  
- acc: 0.9077 - val\_loss: 0.1927 - val\_acc: 0.9421

Epoch 8/20  
60000/60000 [=====] - 11s 185us/step - loss: 0.293  
3 - acc: 0.9113 - val\_loss: 0.1836 - val\_acc: 0.9453

Epoch 9/20  
60000/60000 [=====] - 13s 222us/step - loss: 0.285  
0 - acc: 0.9131 - val\_loss: 0.1797 - val\_acc: 0.9451

Epoch 10/20  
60000/60000 [=====] - 14s 236us/step - loss: 0.271  
5 - acc: 0.9187 - val\_loss: 0.1738 - val\_acc: 0.9465

Epoch 11/20  
60000/60000 [=====] - 8s 141us/step - loss: 0.2611  
- acc: 0.9214 - val\_loss: 0.1671 - val\_acc: 0.9506

Epoch 12/20  
60000/60000 [=====] - 8s 134us/step - loss: 0.2464  
- acc: 0.9252 - val\_loss: 0.1554 - val\_acc: 0.9525

Epoch 13/20  
60000/60000 [=====] - 8s 137us/step - loss: 0.2382  
- acc: 0.9278 - val\_loss: 0.1479 - val\_acc: 0.9554

Epoch 14/20  
60000/60000 [=====] - 8s 136us/step - loss: 0.2275  
- acc: 0.9313 - val\_loss: 0.1375 - val\_acc: 0.9580

Epoch 15/20  
60000/60000 [=====] - 8s 137us/step - loss: 0.2183  
- acc: 0.9337 - val\_loss: 0.1326 - val\_acc: 0.9599

Epoch 16/20  
60000/60000 [=====] - 8s 138us/step - loss: 0.2068  
- acc: 0.9384 - val\_loss: 0.1297 - val\_acc: 0.9613 loss: 0.2066 - ac

Epoch 17/20  
60000/60000 [=====] - 8s 139us/step - loss: 0.2011  
- acc: 0.9395 - val\_loss: 0.1181 - val\_acc: 0.9646

Epoch 18/20  
60000/60000 [=====] - 8s 137us/step - loss: 0.1886  
- acc: 0.9435 - val\_loss: 0.1145 - val\_acc: 0.9658

Epoch 19/20  
60000/60000 [=====] - 8s 138us/step - loss: 0.1821

- acc: 0.9451 - val\_loss: 0.1104 - val\_acc: 0.9662

Epoch 20/20

60000/60000 [=====] - 8s 139us/step - loss: 0.1739

- acc: 0.9473 - val\_loss: 0.1093 - val\_acc: 0.9679



```
In [0]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=
nb_epoch, verbose=1, validation_data=(X_test, Y_test))

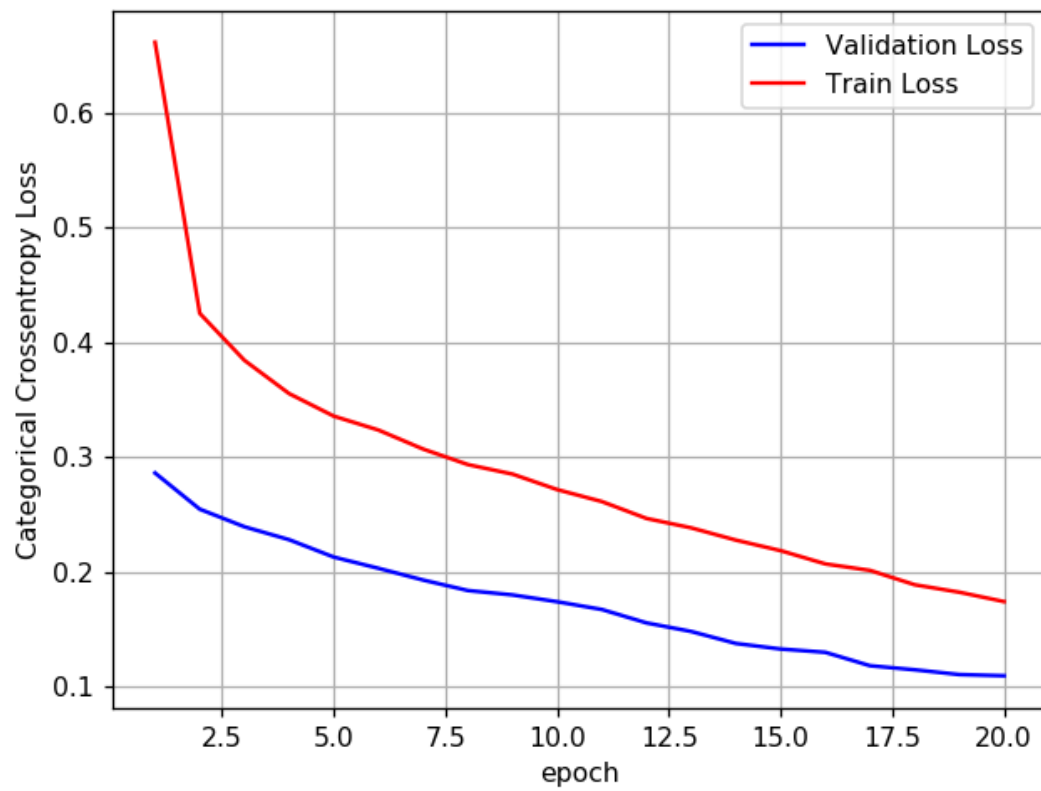
# we will get val_loss and val_acc only when you pass the paramter validati
on_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1093290721397847

Test accuracy: 0.9679



```

In [0]: w_after = model_drop.get_weights()

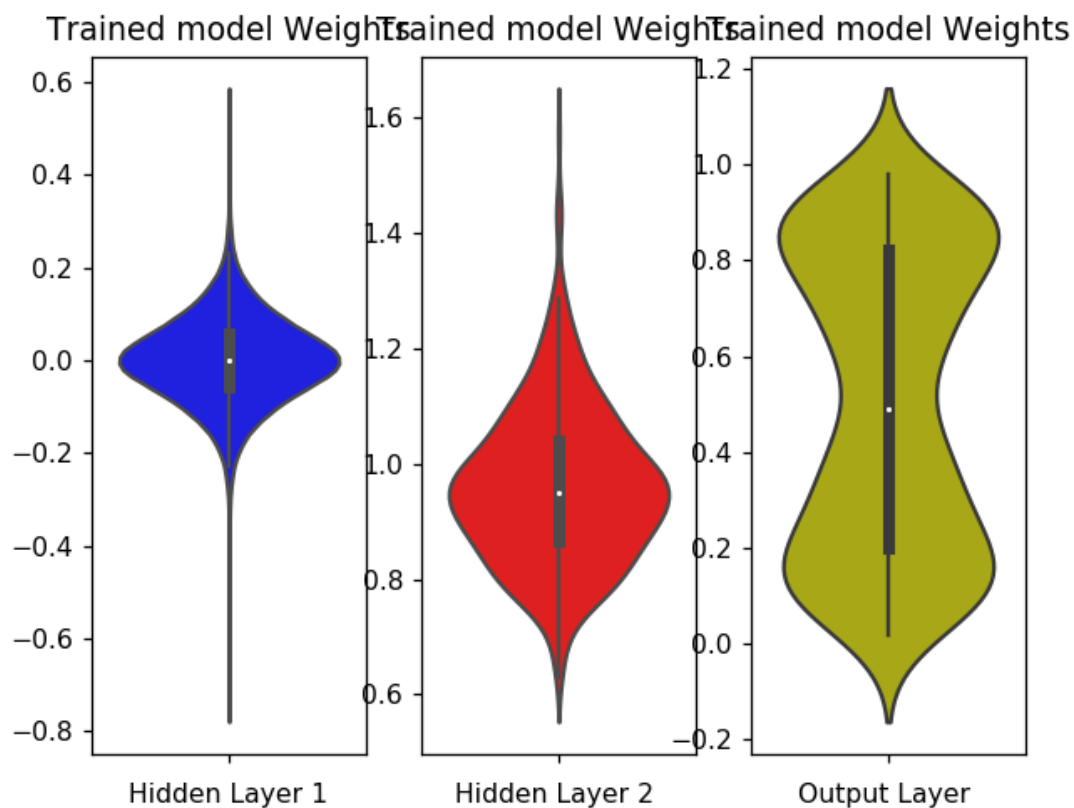
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## Hyper-parameter tuning of Keras models using Sklearn

```
In [0]: from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel
_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(
mean=0.0, stddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], op
timizer='adam')

    return model
```

```
In [0]: # https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid','relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, bat
ch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
In [0]: print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_para
ms_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

Best: 0.975633 using {'activ': 'relu'}
0.974650 (0.001138) with: {'activ': 'sigmoid'}
0.975633 (0.002812) with: {'activ': 'relu'}
```

## 2-layer MLP + Relu + adam

No Dropout and Batch normalization

```
In [0]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_md1_res(x_val, trn_los, tst_los, tst_scr, tst_acc):
    # Visualize loss history
    plt.figure(figsize=(16,16))
    plt.plot(x_val, trn_los, 'r--')
    plt.plot(x_val, tst_los, 'b-')
    plt.legend(['Training Loss', 'Test Loss'])
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.grid()
    plt.show();

    print('Test score:', tst_scr)
    print('Test accuracy:', tst_acc)
```

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20
```

```
In [0]: #MLP
mdl_relu = Sequential()
mdl_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=42)))
mdl_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.150, seed=42)))
mdl_relu.add(Dense(output_dim, activation='softmax'))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:66: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

```
In [0]: mdl_relu.summary()  
print(y_train.shape, y_test.shape)
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dense_2 (Dense)	(None, 128)	65664
dense_3 (Dense)	(None, 10)	1290

=====  
Total params: 468,874  
Trainable params: 468,874  
Non-trainable params: 0  
=====  
(60000, 10) (10000, 10)

```
In [0]: #set optimizer and loss  
mdl_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=  
['accuracy'])
```

```
In [0]: history = mdl_relu.fit(X_train,y_train,batch_size=batch_size,epochs=no_epochs,verbose=1,validation_data=(x_test,y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20  
60000/60000 [=====] - 3s 55us/step - loss: 0.0046  
- acc: 0.9987 - val\_loss: 0.1365 - val\_acc: 0.9821

Epoch 2/20  
60000/60000 [=====] - 3s 50us/step - loss: 0.0069  
- acc: 0.9981 - val\_loss: 0.1310 - val\_acc: 0.9809

Epoch 3/20  
60000/60000 [=====] - 3s 50us/step - loss: 0.0034  
- acc: 0.9989 - val\_loss: 0.1246 - val\_acc: 0.9818

Epoch 4/20  
60000/60000 [=====] - 3s 51us/step - loss: 0.0021  
- acc: 0.9993 - val\_loss: 0.1341 - val\_acc: 0.9814

Epoch 5/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0036  
- acc: 0.9990 - val\_loss: 0.1272 - val\_acc: 0.9814

Epoch 6/20  
60000/60000 [=====] - 3s 51us/step - loss: 0.0059  
- acc: 0.9985 - val\_loss: 0.1455 - val\_acc: 0.9788

Epoch 7/20  
60000/60000 [=====] - 3s 51us/step - loss: 0.0054  
- acc: 0.9986 - val\_loss: 0.1321 - val\_acc: 0.9821

Epoch 8/20  
60000/60000 [=====] - 3s 51us/step - loss: 0.0055  
- acc: 0.9985 - val\_loss: 0.1452 - val\_acc: 0.9799

Epoch 9/20  
60000/60000 [=====] - 3s 51us/step - loss: 0.0033  
- acc: 0.9990 - val\_loss: 0.1300 - val\_acc: 0.9821

Epoch 10/20  
60000/60000 [=====] - 3s 51us/step - loss: 0.0033  
- acc: 0.9992 - val\_loss: 0.1189 - val\_acc: 0.9833

Epoch 11/20  
60000/60000 [=====] - 3s 52us/step - loss: 5.0557e-04  
- acc: 0.9999 - val\_loss: 0.1087 - val\_acc: 0.9840

Epoch 12/20  
60000/60000 [=====] - 3s 50us/step - loss: 0.0058  
- acc: 0.9985 - val\_loss: 0.1919 - val\_acc: 0.9756

Epoch 13/20  
60000/60000 [=====] - 3s 51us/step - loss: 0.0073  
- acc: 0.9983 - val\_loss: 0.1345 - val\_acc: 0.9820

Epoch 14/20  
60000/60000 [=====] - 3s 50us/step - loss: 0.0025  
- acc: 0.9993 - val\_loss: 0.1193 - val\_acc: 0.9834

Epoch 15/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0028  
- acc: 0.9992 - val\_loss: 0.1254 - val\_acc: 0.9825

Epoch 16/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0035  
- acc: 0.9991 - val\_loss: 0.1440 - val\_acc: 0.9816

Epoch 17/20  
60000/60000 [=====] - 3s 51us/step - loss: 0.0016  
- acc: 0.9995 - val\_loss: 0.1338 - val\_acc: 0.9820

Epoch 18/20  
60000/60000 [=====] - 3s 50us/step - loss: 0.0012  
- acc: 0.9996 - val\_loss: 0.1444 - val\_acc: 0.9820

Epoch 19/20  
60000/60000 [=====] - 3s 49us/step - loss: 0.0084



- acc: 0.9981 - val\_loss: 0.1425 - val\_acc: 0.9816

Epoch 20/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0027

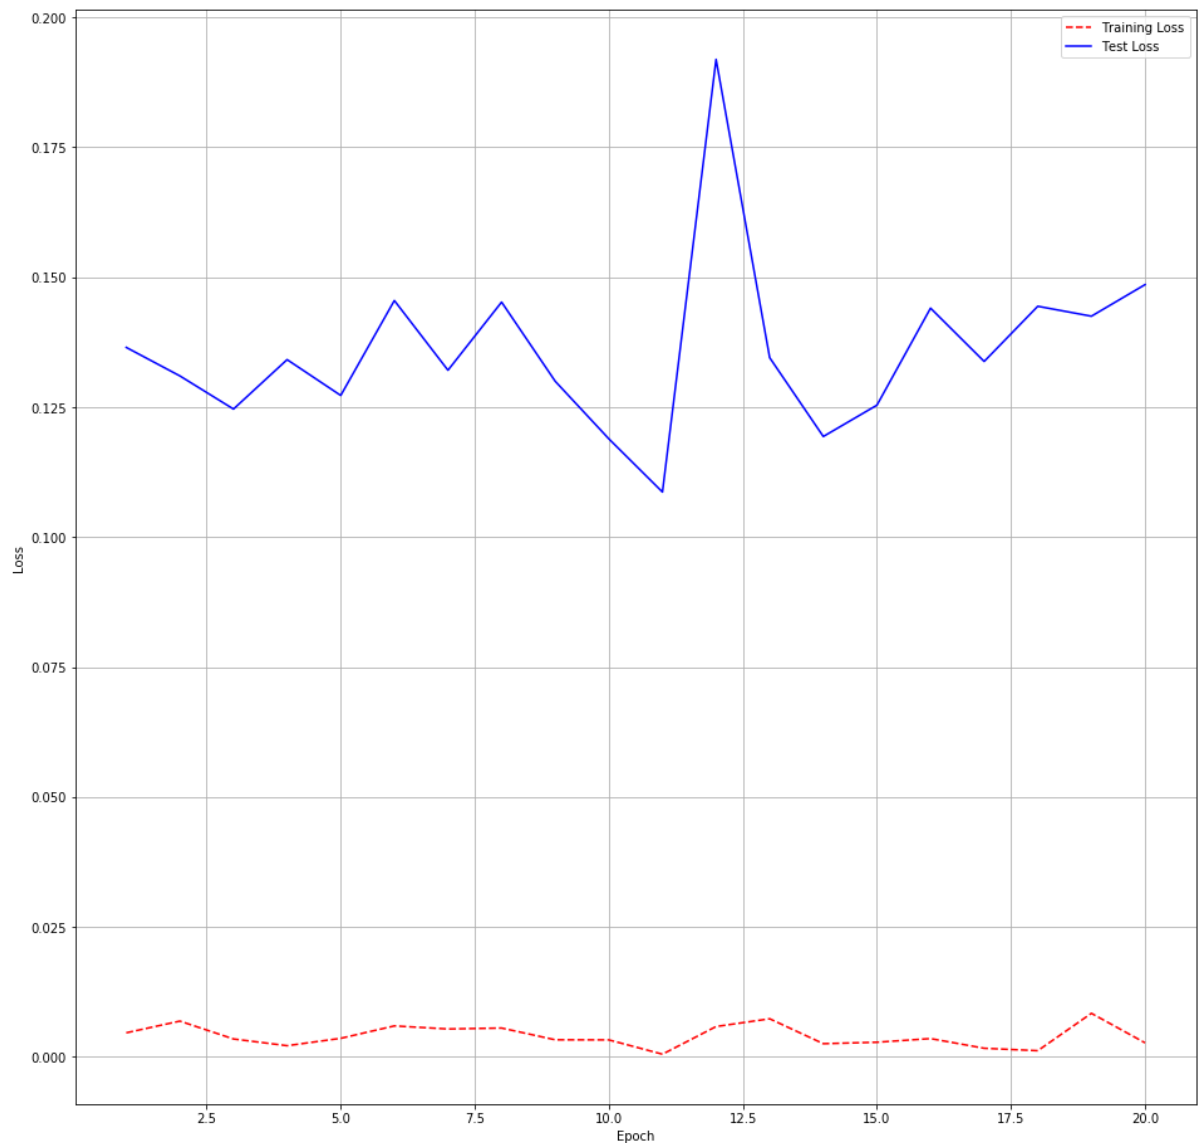
- acc: 0.9993 - val\_loss: 0.1486 - val\_acc: 0.9796

```
In [0]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history.history['loss']
test_loss = history.history['val_loss']

# Create count of the number of epochs
score = mdl_relu.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.14856573140081578

Test accuracy: 0.9796

## 2-Layer MLP + ReLu + Adam + Dropout

### No Batch Normalization

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20
```

```
In [0]: #MLP
mdl_relu_2 = Sequential()
mdl_relu_2.add(Dense(512, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_2.add(Dropout(rate=0.5))
mdl_relu_2.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_2.add(Dropout(rate=0.2))
mdl_relu_2.add(Dense(output_dim,activation='softmax'))
```

```
In [0]: mdl_relu_2.summary()
print(y_train.shape, y_test.shape)
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_10 (Dense)	(None, 512)	401920
dropout_5 (Dropout)	(None, 512)	0
dense_11 (Dense)	(None, 128)	65664
dropout_6 (Dropout)	(None, 128)	0
dense_12 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		
=====	=====	=====
(60000, 10)	(10000, 10)	

```
In [0]: #set optimizer and loss
mdl_relu_2.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [0]: history_2 = mdl_relu_2.fit(X_train,y_train,batch_size=batch_size,epochs=no_
epoch,verbose=1,validation_data=(x_test,y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20  
60000/60000 [=====] - 4s 63us/step - loss: 0.6340  
- acc: 0.8135 - val\_loss: 0.1636 - val\_acc: 0.9517

Epoch 2/20  
60000/60000 [=====] - 3s 55us/step - loss: 0.2473  
- acc: 0.9236 - val\_loss: 0.1268 - val\_acc: 0.9604

Epoch 3/20  
60000/60000 [=====] - 3s 54us/step - loss: 0.1860  
- acc: 0.9437 - val\_loss: 0.1047 - val\_acc: 0.9671

Epoch 4/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.1536  
- acc: 0.9528 - val\_loss: 0.0907 - val\_acc: 0.9712

Epoch 5/20  
60000/60000 [=====] - 3s 54us/step - loss: 0.1290  
- acc: 0.9602 - val\_loss: 0.0871 - val\_acc: 0.9737

Epoch 6/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.1157  
- acc: 0.9640 - val\_loss: 0.0765 - val\_acc: 0.9759

Epoch 7/20  
60000/60000 [=====] - 3s 55us/step - loss: 0.1036  
- acc: 0.9676 - val\_loss: 0.0737 - val\_acc: 0.9774

Epoch 8/20  
60000/60000 [=====] - 3s 55us/step - loss: 0.0944  
- acc: 0.9699 - val\_loss: 0.0731 - val\_acc: 0.9789

Epoch 9/20  
60000/60000 [=====] - 3s 55us/step - loss: 0.0839  
- acc: 0.9732 - val\_loss: 0.0716 - val\_acc: 0.9800

Epoch 10/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0811  
- acc: 0.9746 - val\_loss: 0.0692 - val\_acc: 0.9788

Epoch 11/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0730  
- acc: 0.9764 - val\_loss: 0.0688 - val\_acc: 0.9798

Epoch 12/20  
60000/60000 [=====] - 3s 54us/step - loss: 0.0696  
- acc: 0.9781 - val\_loss: 0.0660 - val\_acc: 0.9807

Epoch 13/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0645  
- acc: 0.9790 - val\_loss: 0.0628 - val\_acc: 0.9814

Epoch 14/20  
60000/60000 [=====] - 3s 54us/step - loss: 0.0604  
- acc: 0.9810 - val\_loss: 0.0616 - val\_acc: 0.9822

Epoch 15/20  
60000/60000 [=====] - 3s 55us/step - loss: 0.0589  
- acc: 0.9807 - val\_loss: 0.0622 - val\_acc: 0.9826

Epoch 16/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0560  
- acc: 0.9819 - val\_loss: 0.0619 - val\_acc: 0.9820

Epoch 17/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0491  
- acc: 0.9843 - val\_loss: 0.0663 - val\_acc: 0.9822

Epoch 18/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0508  
- acc: 0.9836 - val\_loss: 0.0649 - val\_acc: 0.9820

Epoch 19/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0502

- acc: 0.9839 - val\_loss: 0.0613 - val\_acc: 0.9832

Epoch 20/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0462

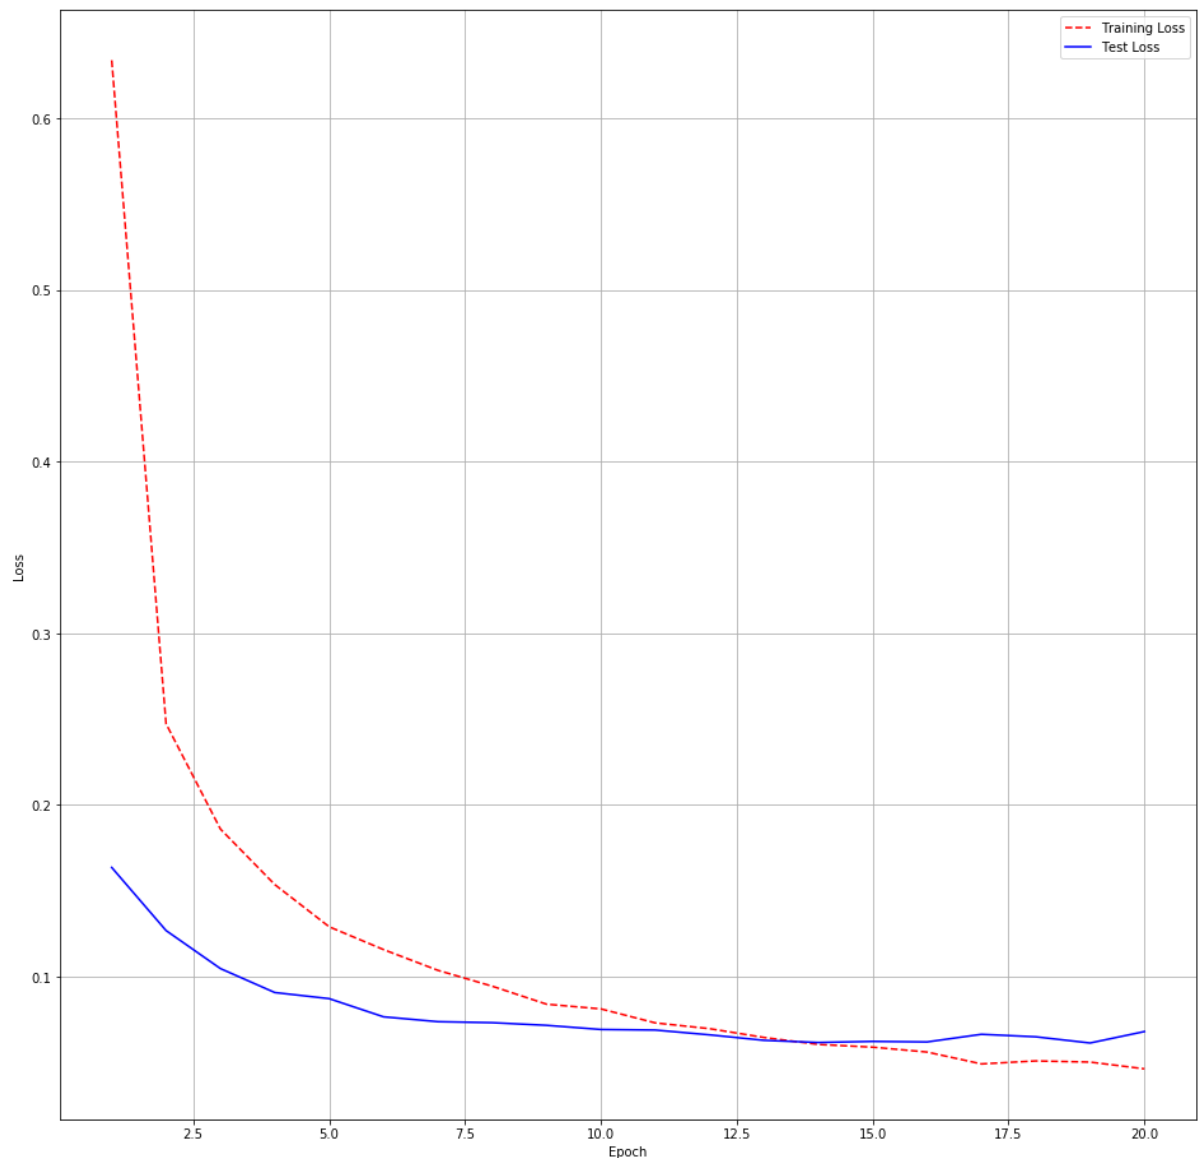
- acc: 0.9852 - val\_loss: 0.0680 - val\_acc: 0.9816

```
In [0]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_2.history['loss']
test_loss = history_2.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_2.evaluate(x_test, y_test, verbose=0)
plt_mdl_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.06795536802747083

Test accuracy: 0.9816

## **2-Layer MLP + ReLu + Adam + Dropout + Batch Normalization**

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20

#MLP
mdl_relu_3 = Sequential()
mdl_relu_3.add(Dense(512, activation='relu',input_shape=(input_dim,),kernel
_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_3.add(Dropout(rate=0.5))
mdl_relu_3.add(BatchNormalization())
mdl_relu_3.add(Dense(128, activation='relu',kernel_initializer=RandomNormal
(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_3.add(Dropout(rate=0.2))
mdl_relu_3.add(BatchNormalization())
mdl_relu_3.add(Dense(output_dim,activation='softmax'))

mdl_relu_3.summary()
mdl_relu_3.compile(optimizer='adam',loss='categorical_crossentropy', metric
s=['accuracy'])
history_3 = mdl_relu_3.fit(X_train,y_train,batch_size=batch_size,epochs=no_
epoch,verbose=1,validation_data=(x_test,y_test))
```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
dense_25 (Dense)	(None, 512)	401920
dropout_15 (Dropout)	(None, 512)	0
batch_normalization_7 (Batch Normalization)	(None, 512)	2048
dense_26 (Dense)	(None, 128)	65664
dropout_16 (Dropout)	(None, 128)	0
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dense_27 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 106us/step - loss: 0.5043  
- acc: 0.8422 - val\_loss: 0.1645 - val\_acc: 0.9514

Epoch 2/20

60000/60000 [=====] - 5s 86us/step - loss: 0.2239  
- acc: 0.9333 - val\_loss: 0.1166 - val\_acc: 0.9645

Epoch 3/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1709  
- acc: 0.9476 - val\_loss: 0.0951 - val\_acc: 0.9696

Epoch 4/20

60000/60000 [=====] - 5s 87us/step - loss: 0.1390  
- acc: 0.9577 - val\_loss: 0.0830 - val\_acc: 0.9743

Epoch 5/20

60000/60000 [=====] - 5s 90us/step - loss: 0.1246  
- acc: 0.9616 - val\_loss: 0.0769 - val\_acc: 0.9757

Epoch 6/20

60000/60000 [=====] - 5s 89us/step - loss: 0.1081  
- acc: 0.9663 - val\_loss: 0.0716 - val\_acc: 0.9772

Epoch 7/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0996  
- acc: 0.9687 - val\_loss: 0.0690 - val\_acc: 0.9788

Epoch 8/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0889  
- acc: 0.9719 - val\_loss: 0.0658 - val\_acc: 0.9798

Epoch 9/20

60000/60000 [=====] - 5s 86us/step - loss: 0.0841  
- acc: 0.9733 - val\_loss: 0.0595 - val\_acc: 0.9821

Epoch 10/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0760  
- acc: 0.9752 - val\_loss: 0.0597 - val\_acc: 0.9812

Epoch 11/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0788  
- acc: 0.9747 - val\_loss: 0.0597 - val\_acc: 0.9818

Epoch 12/20



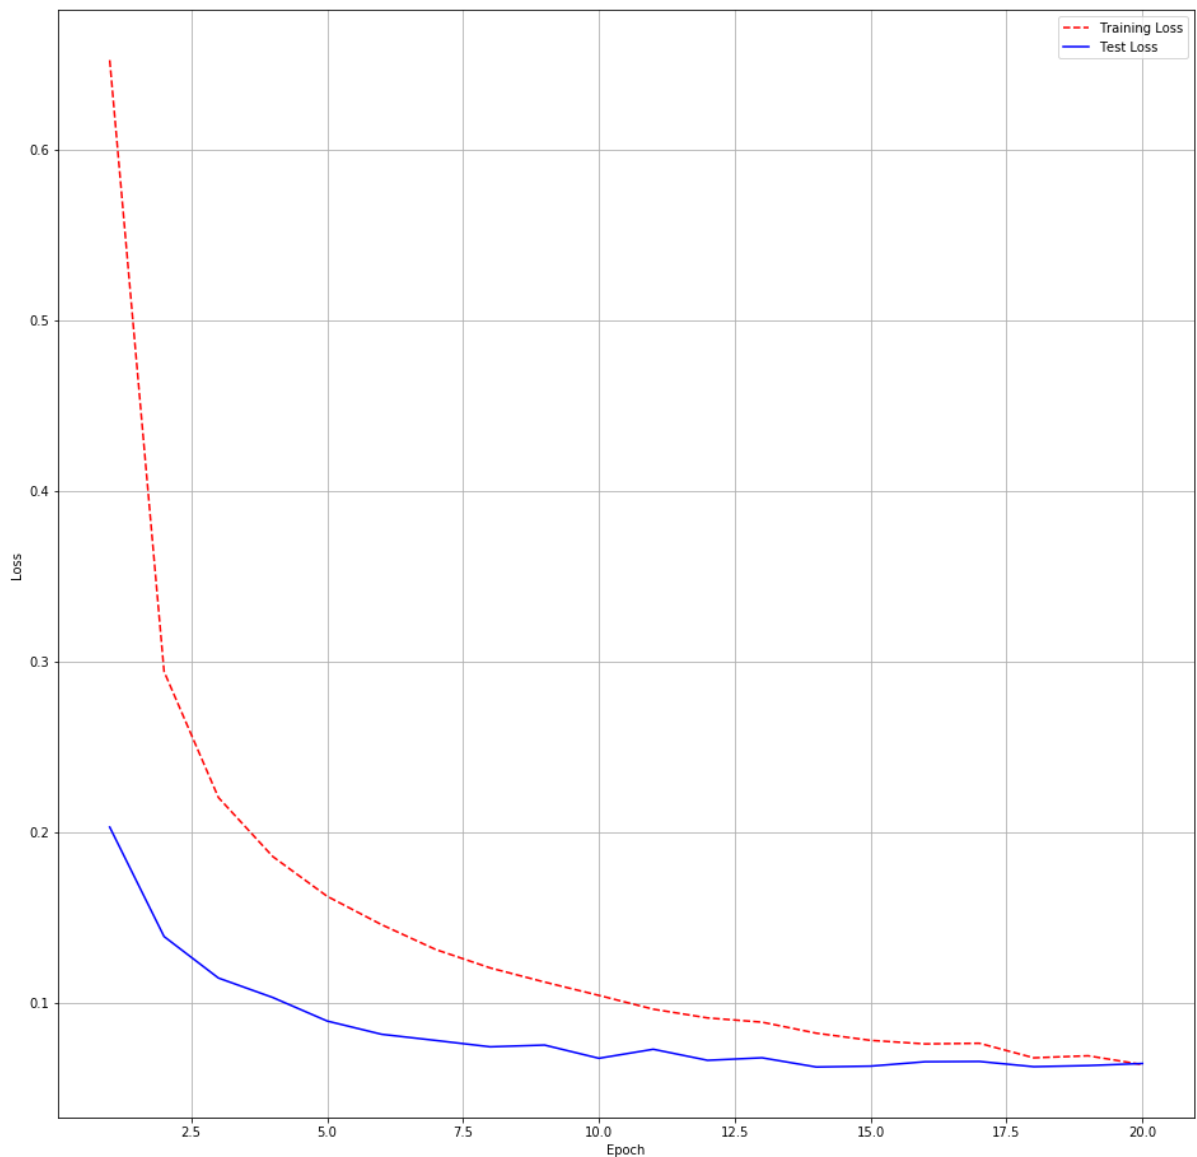
```
60000/60000 [=====] - 5s 87us/step - loss: 0.0722
- acc: 0.9765 - val_loss: 0.0592 - val_acc: 0.9814
Epoch 13/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0667
- acc: 0.9789 - val_loss: 0.0584 - val_acc: 0.9824
Epoch 14/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0623
- acc: 0.9795 - val_loss: 0.0585 - val_acc: 0.9827
Epoch 15/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0608
- acc: 0.9805 - val_loss: 0.0529 - val_acc: 0.9847
Epoch 16/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0555
- acc: 0.9815 - val_loss: 0.0559 - val_acc: 0.9825
Epoch 17/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0569
- acc: 0.9815 - val_loss: 0.0545 - val_acc: 0.9825
Epoch 18/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0510
- acc: 0.9831 - val_loss: 0.0544 - val_acc: 0.9842
Epoch 19/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0503
- acc: 0.9836 - val_loss: 0.0559 - val_acc: 0.9819
Epoch 20/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0504
- acc: 0.9838 - val_loss: 0.0561 - val_acc: 0.9831
```

```
In [0]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_3.history['loss']
test_loss = history_3.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_3.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.06436893670770805

Test accuracy: 0.9809

# **2-Layer MLP + Adam + Relu + Batch Normalization**

**No Dropout**

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20

#MLP
mdl_relu_4 = Sequential()
mdl_relu_4.add(Dense(512, activation='relu',input_shape=(input_dim,),kernel
_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_4.add(BatchNormalization())
mdl_relu_4.add(Dense(128, activation='relu',kernel_initializer=RandomNormal
(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_4.add(BatchNormalization())
mdl_relu_4.add(Dense(output_dim,activation='softmax'))

mdl_relu_4.summary()
mdl_relu_4.compile(optimizer='adam',loss='categorical_crossentropy', metric
s=['accuracy'])
history_4 = mdl_relu_4.fit(X_train,y_train,batch_size=batch_size,epochs=no_
epoch,verbose=1,validation_data=(x_test,y_test))
```

Model: "sequential\_12"

Layer (type)	Output Shape	Param #
dense_34 (Dense)	(None, 512)	401920
batch_normalization_13 (Batch Normalization)	(None, 512)	2048
dense_35 (Dense)	(None, 128)	65664
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dense_36 (Dense)	(None, 10)	1290

Total params: 471,434  
 Trainable params: 470,154  
 Non-trainable params: 1,280

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 110us/step - loss: 0.2537  
 - acc: 0.9250 - val\_loss: 0.1178 - val\_acc: 0.9638

Epoch 2/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0828  
 - acc: 0.9762 - val\_loss: 0.0906 - val\_acc: 0.9737

Epoch 3/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0470  
 - acc: 0.9861 - val\_loss: 0.0820 - val\_acc: 0.9743

Epoch 4/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0285  
 - acc: 0.9922 - val\_loss: 0.0796 - val\_acc: 0.9750

Epoch 5/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0201  
 - acc: 0.9940 - val\_loss: 0.0826 - val\_acc: 0.9767

Epoch 6/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0155  
 - acc: 0.9957 - val\_loss: 0.0802 - val\_acc: 0.9752

Epoch 7/20

60000/60000 [=====] - 5s 84us/step - loss: 0.0136  
 - acc: 0.9960 - val\_loss: 0.0917 - val\_acc: 0.9744

Epoch 8/20

60000/60000 [=====] - 5s 86us/step - loss: 0.0122  
 - acc: 0.9964 - val\_loss: 0.0859 - val\_acc: 0.9770

Epoch 9/20

60000/60000 [=====] - 5s 86us/step - loss: 0.0138  
 - acc: 0.9957 - val\_loss: 0.0833 - val\_acc: 0.9769

Epoch 10/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0091  
 - acc: 0.9971 - val\_loss: 0.0849 - val\_acc: 0.9779

Epoch 11/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0082  
 - acc: 0.9975 - val\_loss: 0.0876 - val\_acc: 0.9792

Epoch 12/20

60000/60000 [=====] - 5s 84us/step - loss: 0.0079  
 - acc: 0.9975 - val\_loss: 0.0791 - val\_acc: 0.9800

Epoch 13/20

60000/60000 [=====] - 5s 84us/step - loss: 0.0070

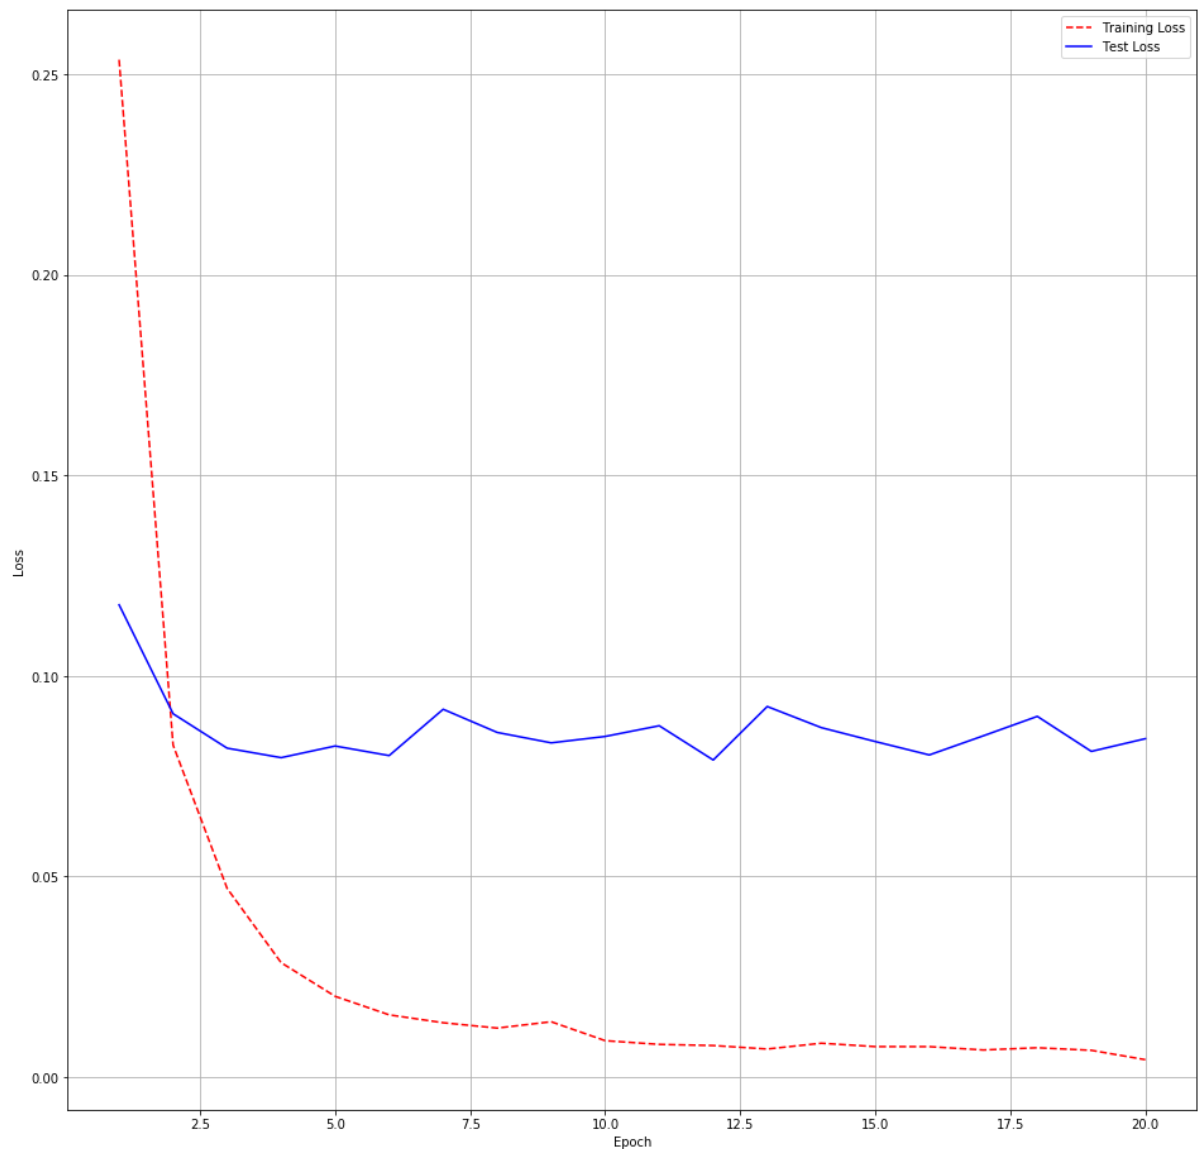
```
- acc: 0.9980 - val_loss: 0.0924 - val_acc: 0.9773
Epoch 14/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0084
- acc: 0.9974 - val_loss: 0.0871 - val_acc: 0.9790
Epoch 15/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0076
- acc: 0.9975 - val_loss: 0.0837 - val_acc: 0.9801
Epoch 16/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0076
- acc: 0.9976 - val_loss: 0.0803 - val_acc: 0.9806
Epoch 17/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0068
- acc: 0.9980 - val_loss: 0.0851 - val_acc: 0.9798
Epoch 18/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0073
- acc: 0.9975 - val_loss: 0.0899 - val_acc: 0.9799
Epoch 19/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0067
- acc: 0.9979 - val_loss: 0.0812 - val_acc: 0.9790
Epoch 20/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0043
- acc: 0.9986 - val_loss: 0.0844 - val_acc: 0.9810
```

```
In [0]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_4.history['loss']
test_loss = history_4.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_4.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.08437914842601021

Test accuracy: 0.981

## [3.1] 3-Layer MLP + ReLu + Adam

### No Dropout and Batch Normalization

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
no_epoch = 20
```



```
In [15]: #MLP
mdl_relu_31 = Sequential()
mdl_relu_31.add(Dense(1024, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_31.add(Dense(512, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_31.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.175, seed=42)))
mdl_relu_31.add(Dense(output_dim,activation='softmax'))
mdl_relu_31.summary()
mdl_relu_31.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_31 = mdl_relu_31.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch,verbose=1,validation_data=(x_test,y_test))
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 1024)	803840
dense_2 (Dense)	(None, 512)	524800
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290
Total params: 1,395,594		
Trainable params: 1,395,594		
Non-trainable params: 0		

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow\_core/python/ops/math\_grad.py:1424: where (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 148us/step - loss: 0.5745  
- acc: 0.9065 - val\_loss: 0.1798 - val\_acc: 0.9523

Epoch 2/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0939  
- acc: 0.9722 - val\_loss: 0.1155 - val\_acc: 0.9648

Epoch 3/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0485  
- acc: 0.9845 - val\_loss: 0.1055 - val\_acc: 0.9713

Epoch 4/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0387  
- acc: 0.9876 - val\_loss: 0.1171 - val\_acc: 0.9689

Epoch 5/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0333  
- acc: 0.9893 - val\_loss: 0.0883 - val\_acc: 0.9761

Epoch 6/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0279  
- acc: 0.9912 - val\_loss: 0.1361 - val\_acc: 0.9656

Epoch 7/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0259  
- acc: 0.9912 - val\_loss: 0.1271 - val\_acc: 0.9665

Epoch 8/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0253  
- acc: 0.9918 - val\_loss: 0.1292 - val\_acc: 0.9685

Epoch 9/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0262  
- acc: 0.9918 - val\_loss: 0.1091 - val\_acc: 0.9736

Epoch 10/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0235  
- acc: 0.9929 - val\_loss: 0.1257 - val\_acc: 0.9724

Epoch 11/20

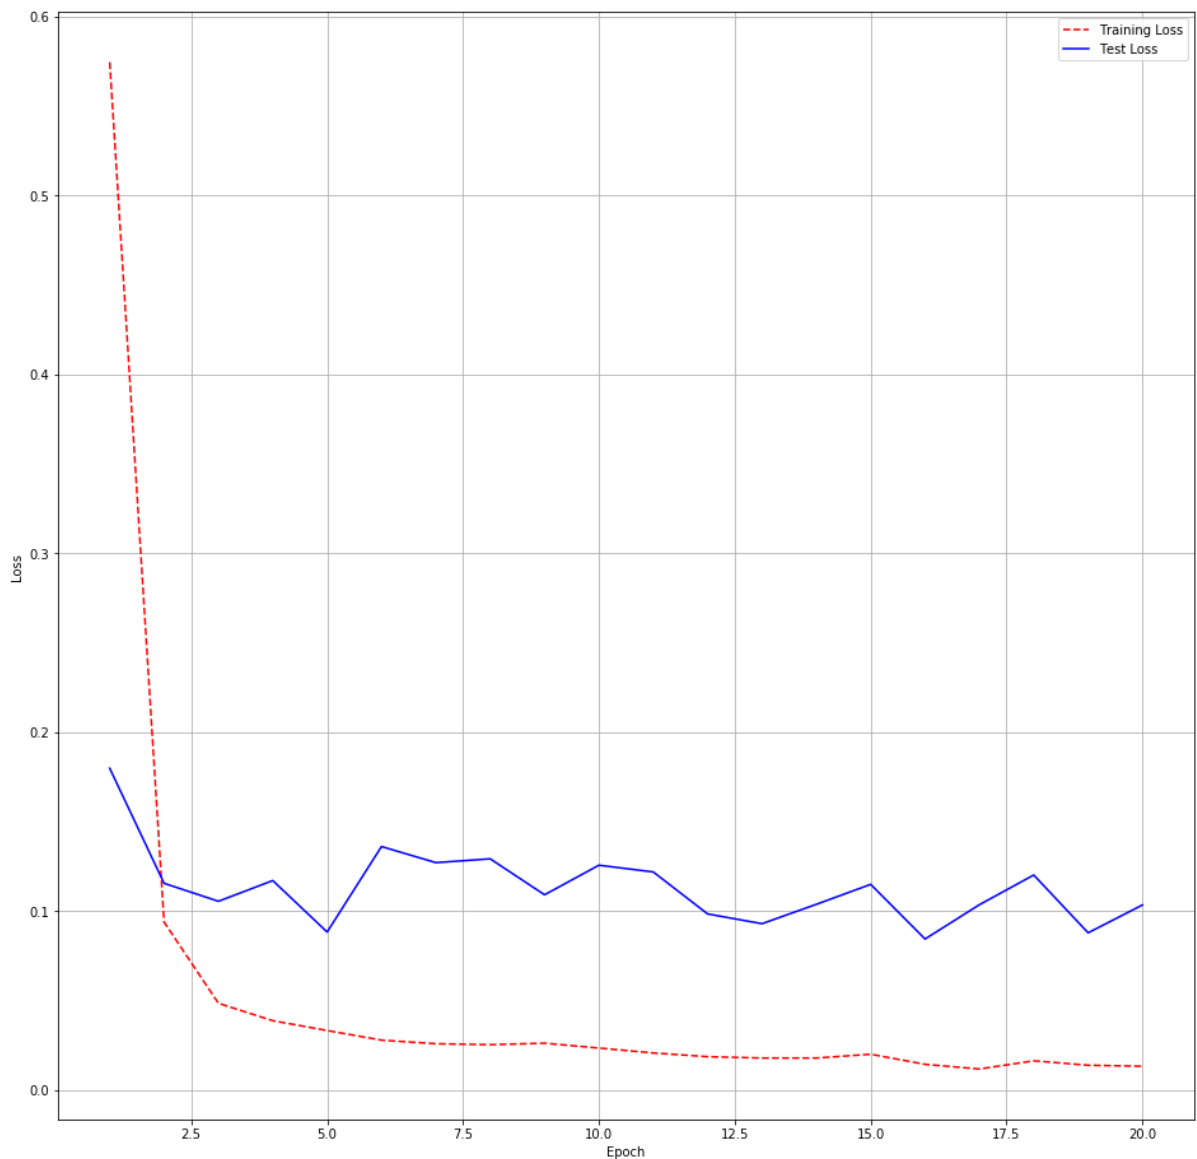
```
60000/60000 [=====] - 4s 68us/step - loss: 0.0207
- acc: 0.9934 - val_loss: 0.1219 - val_acc: 0.9748
Epoch 12/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0186
- acc: 0.9943 - val_loss: 0.0984 - val_acc: 0.9765
Epoch 13/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0179
- acc: 0.9940 - val_loss: 0.0930 - val_acc: 0.9803
Epoch 14/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0179
- acc: 0.9946 - val_loss: 0.1038 - val_acc: 0.9792
Epoch 15/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0200
- acc: 0.9937 - val_loss: 0.1149 - val_acc: 0.9780
Epoch 16/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0144
- acc: 0.9956 - val_loss: 0.0844 - val_acc: 0.9815
Epoch 17/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0118
- acc: 0.9965 - val_loss: 0.1036 - val_acc: 0.9791
Epoch 18/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0163
- acc: 0.9952 - val_loss: 0.1202 - val_acc: 0.9774
Epoch 19/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0139
- acc: 0.9960 - val_loss: 0.0879 - val_acc: 0.9808
Epoch 20/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0133
- acc: 0.9961 - val_loss: 0.1034 - val_acc: 0.9798
```

```
In [19]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_31.history['loss']
test_loss = history_31.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_31.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.10343968682153404

Test accuracy: 0.9798

## **[3.2] 3-Layer MLP + ReLu + Adam + Dropout**

**No Batch Normalization**

```
In [20]: mdl_relu_32 = Sequential()
mdl_relu_32.add(Dense(1024, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_32.add(Dropout(rate=0.5))
mdl_relu_32.add(Dense(512, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_32.add(Dropout(rate=0.5))
mdl_relu_32.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.175, seed=42)))
mdl_relu_32.add(Dropout(rate=0.5))
mdl_relu_32.add(Dense(output_dim,activation='softmax'))
mdl_relu_32.summary()
mdl_relu_32.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_32 = mdl_relu_32.fit(X_train,y_train,batch_size=batch_size,epochs=n_epochs,verbose=1,validation_data=(x_test,y_test))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3733: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 1024)	803840
dropout_1 (Dropout)	(None, 1024)	0
dense_6 (Dense)	(None, 512)	524800
dropout_2 (Dropout)	(None, 512)	0
dense_7 (Dense)	(None, 128)	65664
dropout_3 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 10)	1290

Total params: 1,395,594

Trainable params: 1,395,594

Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 80us/step - loss: 4.2516  
- acc: 0.5804 - val\_loss: 0.3571 - val\_acc: 0.9075

Epoch 2/20

60000/60000 [=====] - 4s 71us/step - loss: 0.7084  
- acc: 0.7998 - val\_loss: 0.2675 - val\_acc: 0.9302

Epoch 3/20

60000/60000 [=====] - 4s 73us/step - loss: 0.4768  
- acc: 0.8635 - val\_loss: 0.2117 - val\_acc: 0.9404

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 0.3735  
- acc: 0.8947 - val\_loss: 0.1791 - val\_acc: 0.9517

Epoch 5/20

60000/60000 [=====] - 4s 73us/step - loss: 0.3039  
- acc: 0.9147 - val\_loss: 0.1552 - val\_acc: 0.9584

Epoch 6/20

60000/60000 [=====] - 4s 73us/step - loss: 0.2615  
- acc: 0.9261 - val\_loss: 0.1374 - val\_acc: 0.9628

Epoch 7/20

60000/60000 [=====] - 4s 73us/step - loss: 0.2271  
- acc: 0.9359 - val\_loss: 0.1238 - val\_acc: 0.9650

Epoch 8/20

60000/60000 [=====] - 4s 70us/step - loss: 0.2051  
- acc: 0.9426 - val\_loss: 0.1159 - val\_acc: 0.9684

Epoch 9/20

60000/60000 [=====] - 4s 75us/step - loss: 0.1847  
- acc: 0.9478 - val\_loss: 0.1052 - val\_acc: 0.9696

Epoch 10/20  
60000/60000 [=====] - 4s 73us/step - loss: 0.1678  
- acc: 0.9532 - val\_loss: 0.1019 - val\_acc: 0.9721

Epoch 11/20  
60000/60000 [=====] - 4s 73us/step - loss: 0.1558  
- acc: 0.9569 - val\_loss: 0.0982 - val\_acc: 0.9722

Epoch 12/20  
60000/60000 [=====] - 4s 73us/step - loss: 0.1424  
- acc: 0.9610 - val\_loss: 0.0963 - val\_acc: 0.9716

Epoch 13/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.1330  
- acc: 0.9624 - val\_loss: 0.0888 - val\_acc: 0.9763

Epoch 14/20  
60000/60000 [=====] - 4s 74us/step - loss: 0.1182  
- acc: 0.9667 - val\_loss: 0.0842 - val\_acc: 0.9784

Epoch 15/20  
60000/60000 [=====] - 5s 75us/step - loss: 0.1180  
- acc: 0.9667 - val\_loss: 0.0800 - val\_acc: 0.9784

Epoch 16/20  
60000/60000 [=====] - 4s 72us/step - loss: 0.1079  
- acc: 0.9696 - val\_loss: 0.0832 - val\_acc: 0.9774

Epoch 17/20  
60000/60000 [=====] - 4s 72us/step - loss: 0.1044  
- acc: 0.9701 - val\_loss: 0.0809 - val\_acc: 0.9782

Epoch 18/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.0953  
- acc: 0.9734 - val\_loss: 0.0872 - val\_acc: 0.9789

Epoch 19/20  
60000/60000 [=====] - 4s 72us/step - loss: 0.0919  
- acc: 0.9732 - val\_loss: 0.0833 - val\_acc: 0.9796

Epoch 20/20  
60000/60000 [=====] - 5s 75us/step - loss: 0.0878  
- acc: 0.9755 - val\_loss: 0.0805 - val\_acc: 0.9797

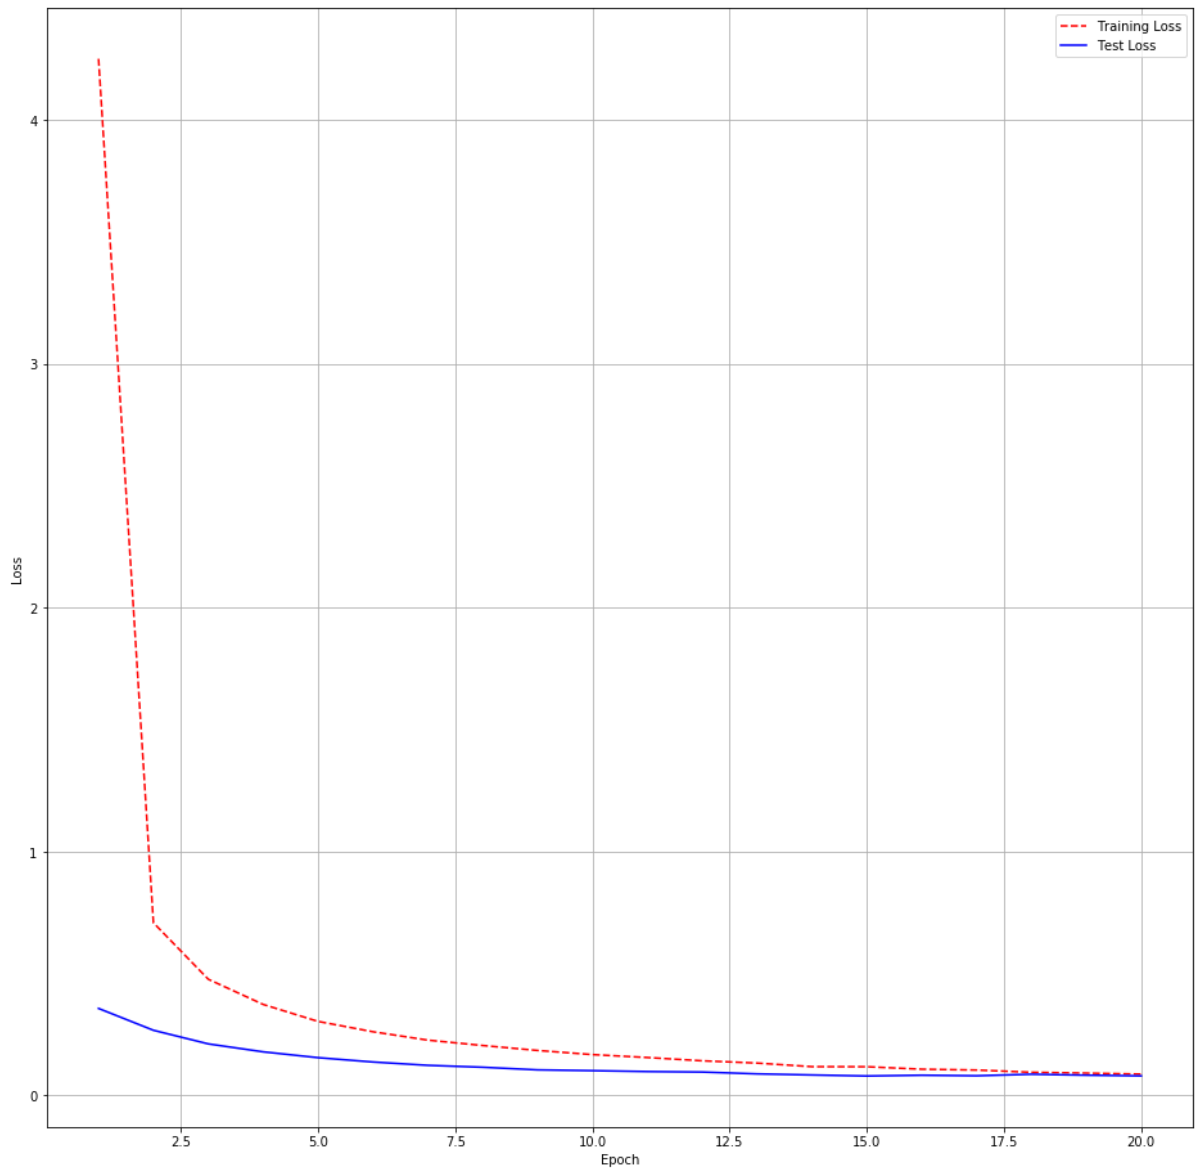


```
In [21]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_32.history['loss']
test_loss = history_32.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_32.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.08047824457662764

Test accuracy: 0.9797

## **[3.3] 3-Layer MLP + ReLu + Adam + Dropout + Batch Normalization**

```
In [22]: mdl_relu_33 = Sequential()
mdl_relu_33.add(Dense(1024, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_33.add(Dropout(rate=0.5))
mdl_relu_33.add(BatchNormalization())
mdl_relu_33.add(Dense(512, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_33.add(Dropout(rate=0.5))
mdl_relu_33.add(BatchNormalization())
mdl_relu_33.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.175, seed=42)))
mdl_relu_33.add(Dropout(rate=0.5))
mdl_relu_33.add(BatchNormalization())
mdl_relu_33.add(Dense(output_dim,activation='softmax'))
mdl_relu_33.summary()
mdl_relu_33.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_33 = mdl_relu_33.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch,verbose=1,validation_data=(x_test,y_test))
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 1024)	803840
dropout_4 (Dropout)	(None, 1024)	0
batch_normalization_1 (Batch Normalization)	(None, 1024)	4096
dense_10 (Dense)	(None, 512)	524800
dropout_5 (Dropout)	(None, 512)	0
batch_normalization_2 (Batch Normalization)	(None, 512)	2048
dense_11 (Dense)	(None, 128)	65664
dropout_6 (Dropout)	(None, 128)	0
batch_normalization_3 (Batch Normalization)	(None, 128)	512
dense_12 (Dense)	(None, 10)	1290
Total params: 1,402,250		
Trainable params: 1,398,922		
Non-trainable params: 3,328		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 141us/step - loss: 0.7473  
- acc: 0.7661 - val\_loss: 0.1916 - val\_acc: 0.9441

Epoch 2/20

60000/60000 [=====] - 8s 127us/step - loss: 0.2936  
- acc: 0.9141 - val\_loss: 0.1262 - val\_acc: 0.9609

Epoch 3/20

60000/60000 [=====] - 7s 124us/step - loss: 0.2149  
- acc: 0.9385 - val\_loss: 0.1017 - val\_acc: 0.9691

Epoch 4/20

60000/60000 [=====] - 7s 123us/step - loss: 0.1701  
- acc: 0.9505 - val\_loss: 0.0907 - val\_acc: 0.9719

Epoch 5/20

60000/60000 [=====] - 7s 124us/step - loss: 0.1497  
- acc: 0.9563 - val\_loss: 0.0824 - val\_acc: 0.9747

Epoch 6/20

60000/60000 [=====] - 7s 118us/step - loss: 0.1292  
- acc: 0.9614 - val\_loss: 0.0704 - val\_acc: 0.9771

Epoch 7/20

60000/60000 [=====] - 8s 125us/step - loss: 0.1177  
- acc: 0.9653 - val\_loss: 0.0706 - val\_acc: 0.9780

Epoch 8/20

60000/60000 [=====] - 7s 124us/step - loss: 0.1094  
- acc: 0.9674 - val\_loss: 0.0642 - val\_acc: 0.9795

Epoch 9/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0991  
- acc: 0.9714 - val\_loss: 0.0699 - val\_acc: 0.9789

Epoch 10/20

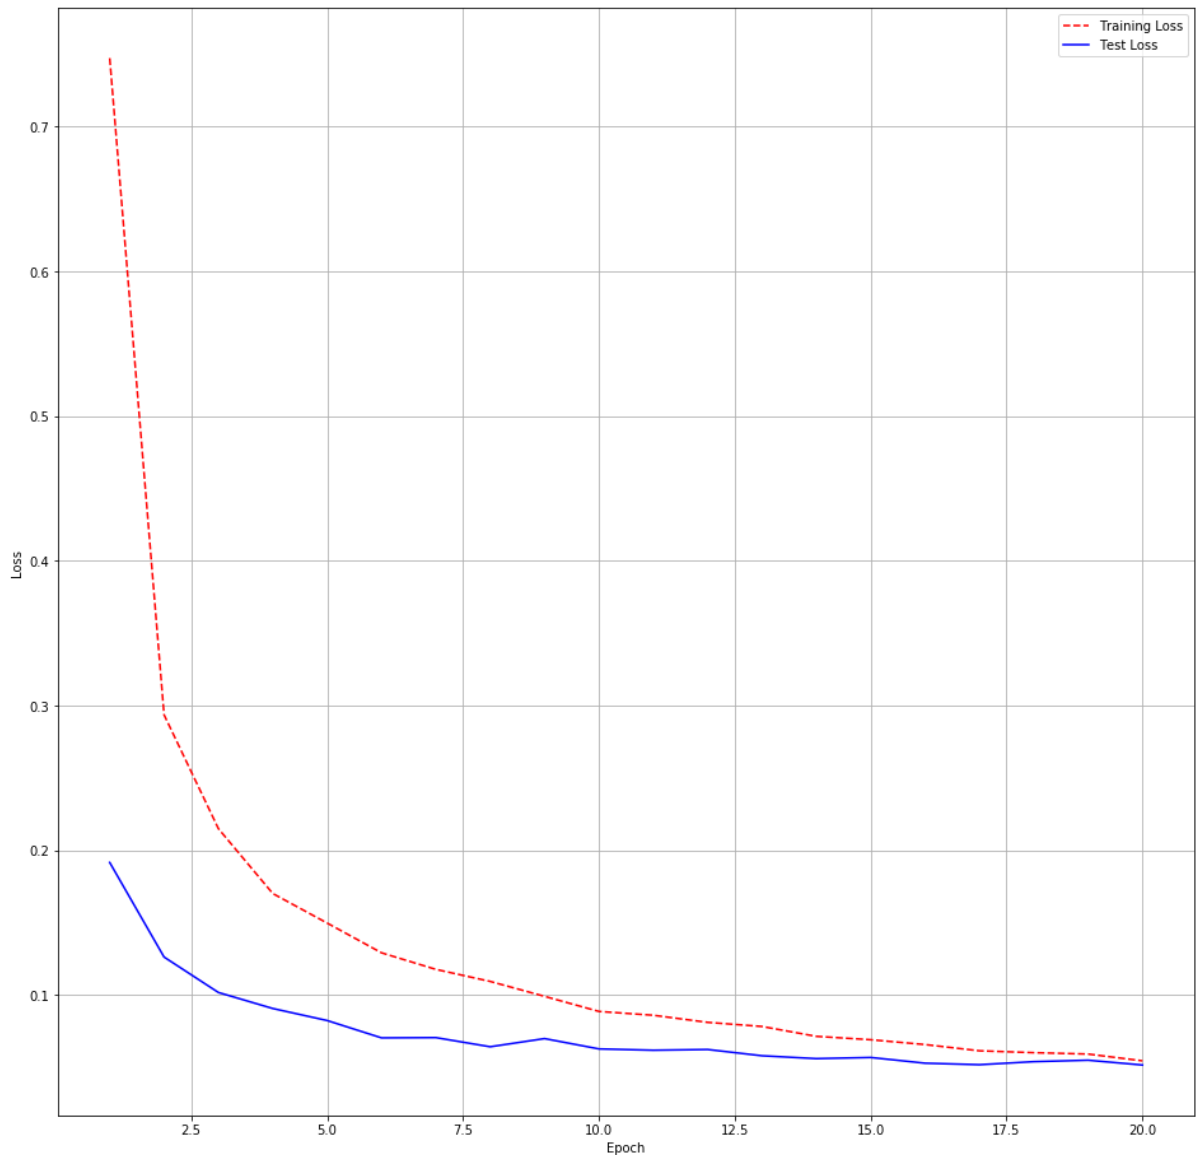
```
60000/60000 [=====] - 8s 127us/step - loss: 0.0886
- acc: 0.9724 - val_loss: 0.0628 - val_acc: 0.9803
Epoch 11/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0860
- acc: 0.9740 - val_loss: 0.0619 - val_acc: 0.9807
Epoch 12/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0811
- acc: 0.9749 - val_loss: 0.0624 - val_acc: 0.9795
Epoch 13/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0783
- acc: 0.9767 - val_loss: 0.0580 - val_acc: 0.9820
Epoch 14/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0714
- acc: 0.9783 - val_loss: 0.0561 - val_acc: 0.9822
Epoch 15/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0691
- acc: 0.9791 - val_loss: 0.0568 - val_acc: 0.9827
Epoch 16/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0658
- acc: 0.9807 - val_loss: 0.0529 - val_acc: 0.9837
Epoch 17/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0615
- acc: 0.9813 - val_loss: 0.0519 - val_acc: 0.9836
Epoch 18/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0602
- acc: 0.9823 - val_loss: 0.0540 - val_acc: 0.9829
Epoch 19/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0592
- acc: 0.9817 - val_loss: 0.0549 - val_acc: 0.9827
Epoch 20/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0545
- acc: 0.9838 - val_loss: 0.0516 - val_acc: 0.9846
```

```
In [23]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_33.history['loss']
test_loss = history_33.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_33.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.05163941752879764

Test accuracy: 0.9846

## **[3.4] 3-Layer MLP + ReLu + Adam + Batch Normalization**

**No Dropout**

```
In [24]: mdl_relu_34 = Sequential()
mdl_relu_34.add(Dense(1024, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_34.add(BatchNormalization())
mdl_relu_34.add(Dense(512, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_34.add(BatchNormalization())
mdl_relu_34.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.175, seed=42)))
mdl_relu_34.add(BatchNormalization())
mdl_relu_34.add(Dense(output_dim,activation='softmax'))
mdl_relu_34.summary()
mdl_relu_34.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_34 = mdl_relu_34.fit(X_train,y_train,batch_size=batch_size,epochs=n_epochs,verbose=1,validation_data=(x_test,y_test))
```



Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 1024)	803840
batch_normalization_4 (Batch Normalization)	(None, 1024)	4096
dense_14 (Dense)	(None, 512)	524800
batch_normalization_5 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_6 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 1,402,250		
Trainable params: 1,398,922		
Non-trainable params: 3,328		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 140us/step - loss: 0.2044  
- acc: 0.9391 - val\_loss: 0.0923 - val\_acc: 0.9718

Epoch 2/20

60000/60000 [=====] - 7s 124us/step - loss: 0.0586  
- acc: 0.9827 - val\_loss: 0.0820 - val\_acc: 0.9730

Epoch 3/20

60000/60000 [=====] - 7s 123us/step - loss: 0.0315  
- acc: 0.9906 - val\_loss: 0.0829 - val\_acc: 0.9768

Epoch 4/20

60000/60000 [=====] - 7s 121us/step - loss: 0.0210  
- acc: 0.9935 - val\_loss: 0.0843 - val\_acc: 0.9752

Epoch 5/20

60000/60000 [=====] - 7s 121us/step - loss: 0.0190  
- acc: 0.9939 - val\_loss: 0.0923 - val\_acc: 0.9737

Epoch 6/20

60000/60000 [=====] - 7s 122us/step - loss: 0.0153  
- acc: 0.9952 - val\_loss: 0.1013 - val\_acc: 0.9718

Epoch 7/20

60000/60000 [=====] - 8s 128us/step - loss: 0.0158  
- acc: 0.9951 - val\_loss: 0.0856 - val\_acc: 0.9753

Epoch 8/20

60000/60000 [=====] - 8s 127us/step - loss: 0.0128  
- acc: 0.9959 - val\_loss: 0.0742 - val\_acc: 0.9809

Epoch 9/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0140  
- acc: 0.9953 - val\_loss: 0.0765 - val\_acc: 0.9793

Epoch 10/20

60000/60000 [=====] - 7s 122us/step - loss: 0.0107  
- acc: 0.9966 - val\_loss: 0.0875 - val\_acc: 0.9779

Epoch 11/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0099  
- acc: 0.9968 - val\_loss: 0.0853 - val\_acc: 0.9790

Epoch 12/20

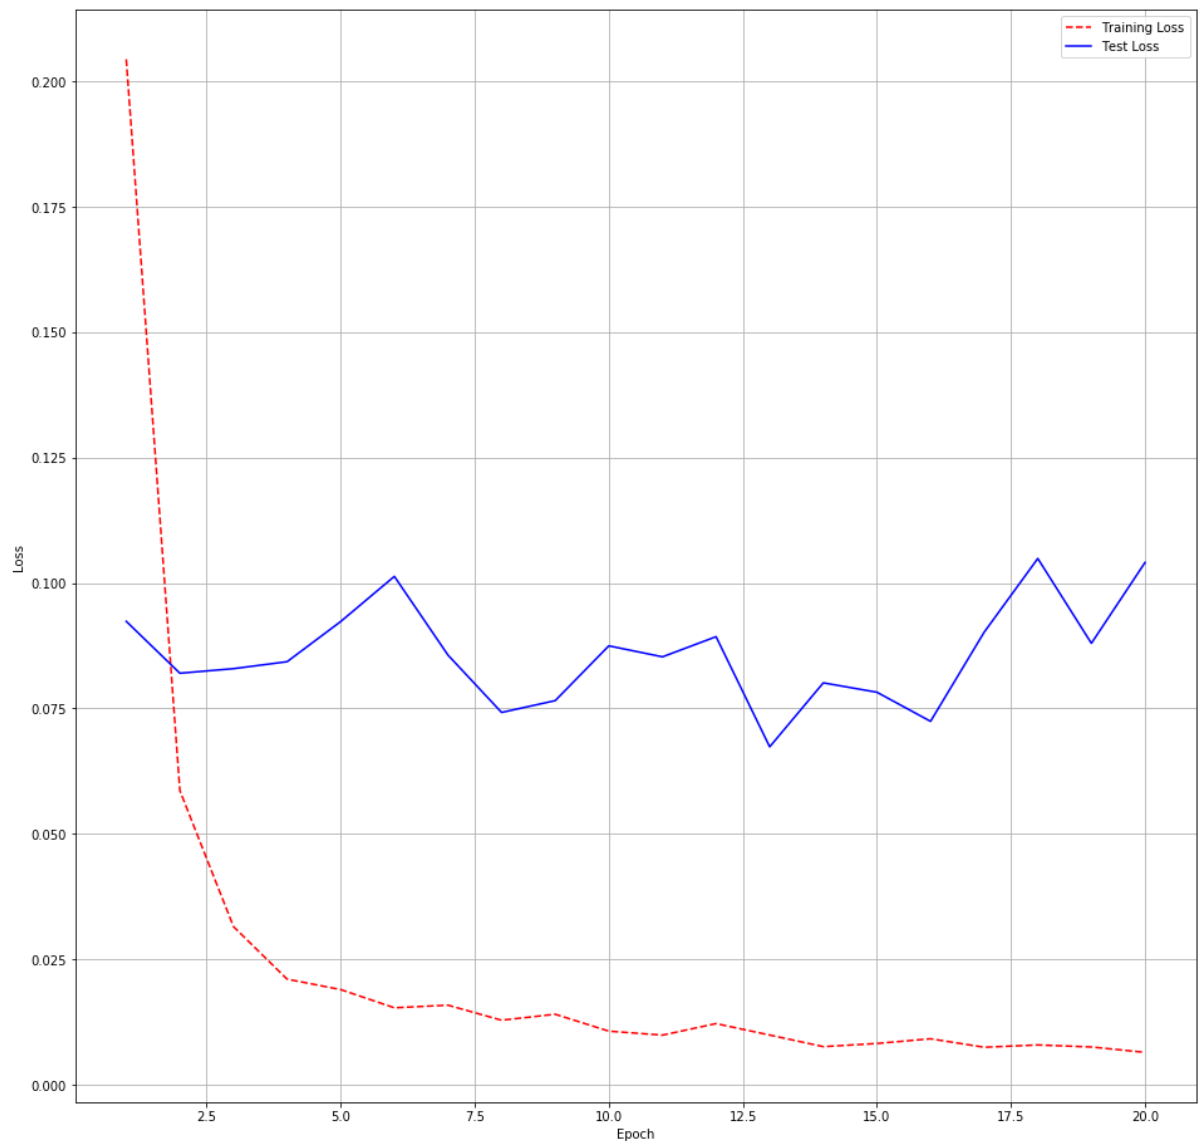
```
60000/60000 [=====] - 7s 124us/step - loss: 0.0122
- acc: 0.9958 - val_loss: 0.0893 - val_acc: 0.9767
Epoch 13/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0099
- acc: 0.9966 - val_loss: 0.0673 - val_acc: 0.9818
Epoch 14/20
60000/60000 [=====] - 8s 125us/step - loss: 0.0076
- acc: 0.9974 - val_loss: 0.0801 - val_acc: 0.9807
Epoch 15/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0082
- acc: 0.9973 - val_loss: 0.0782 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0091
- acc: 0.9969 - val_loss: 0.0724 - val_acc: 0.9817
Epoch 17/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0075
- acc: 0.9976 - val_loss: 0.0902 - val_acc: 0.9772
Epoch 18/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0079
- acc: 0.9973 - val_loss: 0.1049 - val_acc: 0.9773
Epoch 19/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0075
- acc: 0.9976 - val_loss: 0.0880 - val_acc: 0.9788
Epoch 20/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0064
- acc: 0.9977 - val_loss: 0.1041 - val_acc: 0.9771
```

```
In [26]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_34.history['loss']
test_loss = history_34.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_34.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.10407820313923867  
Test accuracy: 0.9771

## **[4.1] 3-Layer MLP + ReLu + Adam**

**No Dropout and Batch Normalization**

```
In [27]: mdl_relu_41 = Sequential()
mdl_relu_41.add(Dense(512, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_41.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_41.add(Dense(64, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.175, seed=42)))
mdl_relu_41.add(Dense(32, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.180, seed=42)))
mdl_relu_41.add(Dense(16, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.190, seed=42)))
mdl_relu_41.add(Dense(output_dim,activation='softmax'))
mdl_relu_41.summary()
mdl_relu_41.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_41 = mdl_relu_41.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch,verbose=1,validation_data=(x_test,y_test))
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 512)	401920
dense_18 (Dense)	(None, 128)	65664
dense_19 (Dense)	(None, 64)	8256
dense_20 (Dense)	(None, 32)	2080
dense_21 (Dense)	(None, 16)	528
dense_22 (Dense)	(None, 10)	170

=====  
 Total params: 478,618  
 Trainable params: 478,618  
 Non-trainable params: 0

---

 Train on 60000 samples, validate on 10000 samples

Epoch 1/20

 60000/60000 [=====] - 5s 79us/step - loss: 0.3879  
 - acc: 0.8846 - val\_loss: 0.1488 - val\_acc: 0.9567

Epoch 2/20

 60000/60000 [=====] - 4s 66us/step - loss: 0.1126  
 - acc: 0.9656 - val\_loss: 0.1091 - val\_acc: 0.9668

Epoch 3/20

 60000/60000 [=====] - 4s 66us/step - loss: 0.0731  
 - acc: 0.9779 - val\_loss: 0.0948 - val\_acc: 0.9720

Epoch 4/20

 60000/60000 [=====] - 4s 66us/step - loss: 0.0487  
 - acc: 0.9850 - val\_loss: 0.0895 - val\_acc: 0.9736

Epoch 5/20

 60000/60000 [=====] - 4s 66us/step - loss: 0.0384  
 - acc: 0.9874 - val\_loss: 0.0966 - val\_acc: 0.9756

Epoch 6/20

 60000/60000 [=====] - 4s 64us/step - loss: 0.0333  
 - acc: 0.9890 - val\_loss: 0.0949 - val\_acc: 0.9732

Epoch 7/20

 60000/60000 [=====] - 4s 65us/step - loss: 0.0236  
 - acc: 0.9921 - val\_loss: 0.1035 - val\_acc: 0.9741

Epoch 8/20

 60000/60000 [=====] - 4s 65us/step - loss: 0.0256  
 - acc: 0.9914 - val\_loss: 0.1030 - val\_acc: 0.9724

Epoch 9/20

 60000/60000 [=====] - 4s 67us/step - loss: 0.0195  
 - acc: 0.9935 - val\_loss: 0.0909 - val\_acc: 0.9781

Epoch 10/20

 60000/60000 [=====] - 4s 68us/step - loss: 0.0160  
 - acc: 0.9948 - val\_loss: 0.1068 - val\_acc: 0.9747

Epoch 11/20

 60000/60000 [=====] - 4s 65us/step - loss: 0.0143  
 - acc: 0.9957 - val\_loss: 0.0936 - val\_acc: 0.9798

Epoch 12/20

 60000/60000 [=====] - 4s 65us/step - loss: 0.0176  
 - acc: 0.9942 - val\_loss: 0.1018 - val\_acc: 0.9758

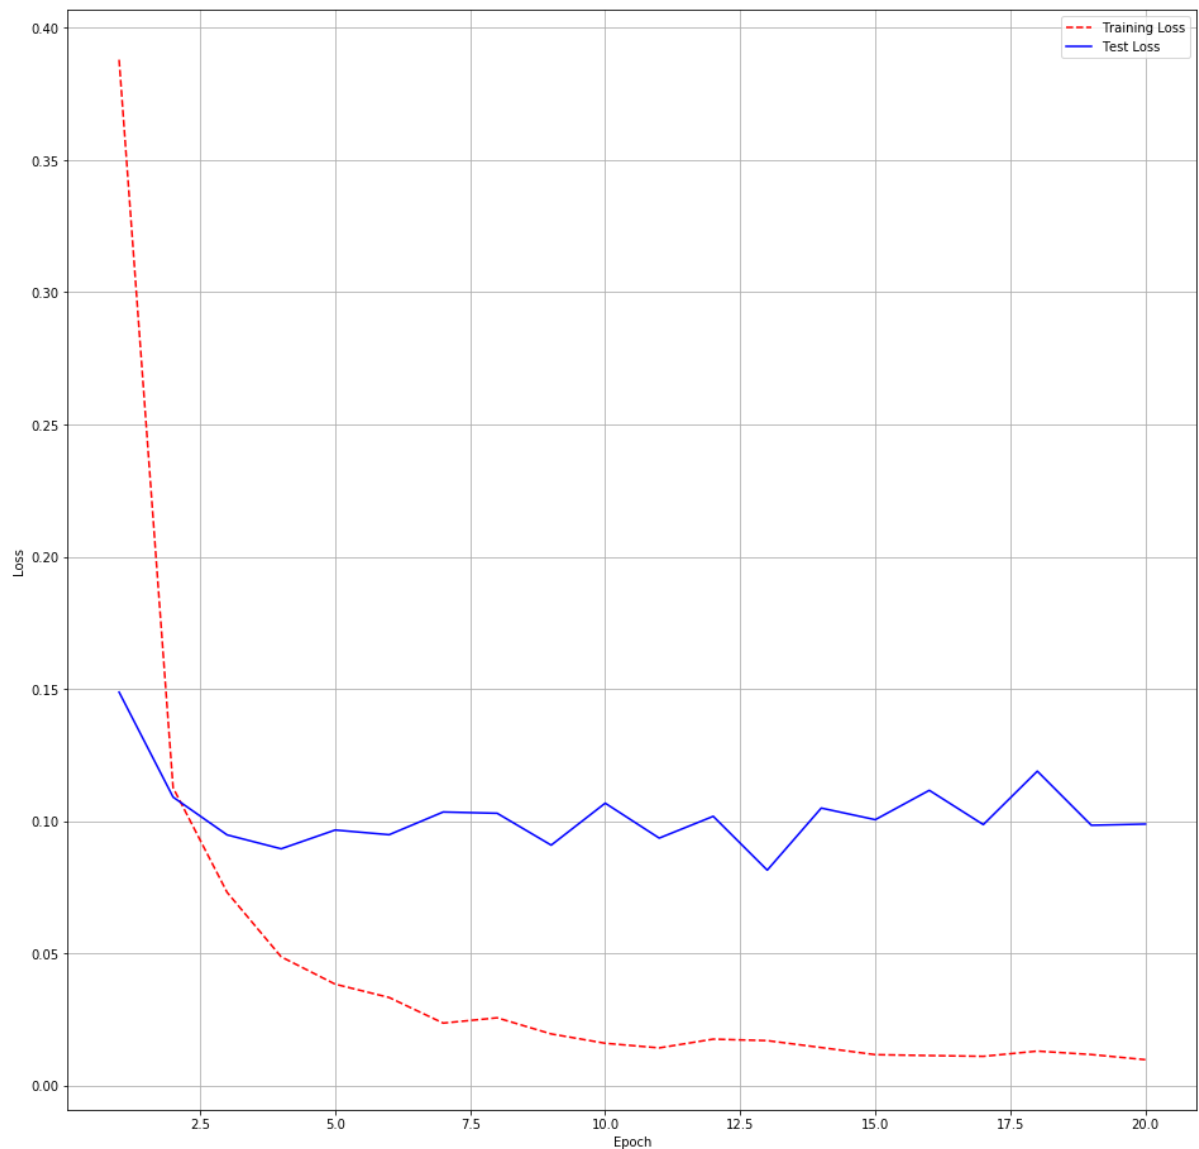
Epoch 13/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0170  
- acc: 0.9947 - val\_loss: 0.0815 - val\_acc: 0.9807  
Epoch 14/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0144  
- acc: 0.9955 - val\_loss: 0.1050 - val\_acc: 0.9758  
Epoch 15/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0117  
- acc: 0.9959 - val\_loss: 0.1006 - val\_acc: 0.9792  
Epoch 16/20  
60000/60000 [=====] - 4s 65us/step - loss: 0.0114  
- acc: 0.9965 - val\_loss: 0.1117 - val\_acc: 0.9774  
Epoch 17/20  
60000/60000 [=====] - 4s 67us/step - loss: 0.0111  
- acc: 0.9962 - val\_loss: 0.0987 - val\_acc: 0.9788  
Epoch 18/20  
60000/60000 [=====] - 4s 67us/step - loss: 0.0130  
- acc: 0.9959 - val\_loss: 0.1189 - val\_acc: 0.9745  
Epoch 19/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0117  
- acc: 0.9962 - val\_loss: 0.0984 - val\_acc: 0.9801  
Epoch 20/20  
60000/60000 [=====] - 4s 64us/step - loss: 0.0098  
- acc: 0.9970 - val\_loss: 0.0989 - val\_acc: 0.9782

```
In [28]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_41.history['loss']
test_loss = history_41.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_41.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.09889793922199387

Test accuracy: 0.9782



## **[4.2] 3-Layer MLP + ReLu + Adam + Dropout**

**No Batch Normalization**

```
In [29]: mdl_relu_42 = Sequential()
mdl_relu_42.add(Dense(512, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_42.add(Dropout(rate=0.5))
mdl_relu_42.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_42.add(Dropout(rate=0.5))
mdl_relu_42.add(Dense(64, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.175, seed=42)))
mdl_relu_42.add(Dropout(rate=0.5))
mdl_relu_42.add(Dense(32, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.180, seed=42)))
mdl_relu_42.add(Dropout(rate=0.5))
mdl_relu_42.add(Dense(16, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.190, seed=42)))
mdl_relu_42.add(Dropout(rate=0.5))
mdl_relu_42.add(Dense(output_dim,activation='softmax'))
mdl_relu_42.summary()
mdl_relu_42.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_42 = mdl_relu_42.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch,verbose=1,validation_data=(x_test,y_test))
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
dense_23 (Dense)	(None, 512)	401920
dropout_7 (Dropout)	(None, 512)	0
dense_24 (Dense)	(None, 128)	65664
dropout_8 (Dropout)	(None, 128)	0
dense_25 (Dense)	(None, 64)	8256
dropout_9 (Dropout)	(None, 64)	0
dense_26 (Dense)	(None, 32)	2080
dropout_10 (Dropout)	(None, 32)	0
dense_27 (Dense)	(None, 16)	528
dropout_11 (Dropout)	(None, 16)	0
dense_28 (Dense)	(None, 10)	170

Total params: 478,618  
 Trainable params: 478,618  
 Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 90us/step - loss: 2.7650  
 - acc: 0.1245 - val\_loss: 2.2475 - val\_acc: 0.1779

Epoch 2/20

60000/60000 [=====] - 4s 71us/step - loss: 2.1208  
 - acc: 0.2122 - val\_loss: 1.7012 - val\_acc: 0.3865

Epoch 3/20

60000/60000 [=====] - 4s 72us/step - loss: 1.7758  
 - acc: 0.3201 - val\_loss: 1.4016 - val\_acc: 0.4645

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 1.5281  
 - acc: 0.4101 - val\_loss: 1.2552 - val\_acc: 0.5403

Epoch 5/20

60000/60000 [=====] - 4s 72us/step - loss: 1.3543  
 - acc: 0.4754 - val\_loss: 1.1145 - val\_acc: 0.5795

Epoch 6/20

60000/60000 [=====] - 4s 73us/step - loss: 1.2332  
 - acc: 0.5396 - val\_loss: 0.9837 - val\_acc: 0.6578

Epoch 7/20

60000/60000 [=====] - 4s 71us/step - loss: 1.1352  
 - acc: 0.5759 - val\_loss: 0.8884 - val\_acc: 0.6652

Epoch 8/20

60000/60000 [=====] - 4s 72us/step - loss: 1.0563  
 - acc: 0.6120 - val\_loss: 0.8219 - val\_acc: 0.7128

Epoch 9/20

60000/60000 [=====] - 4s 69us/step - loss: 0.9792

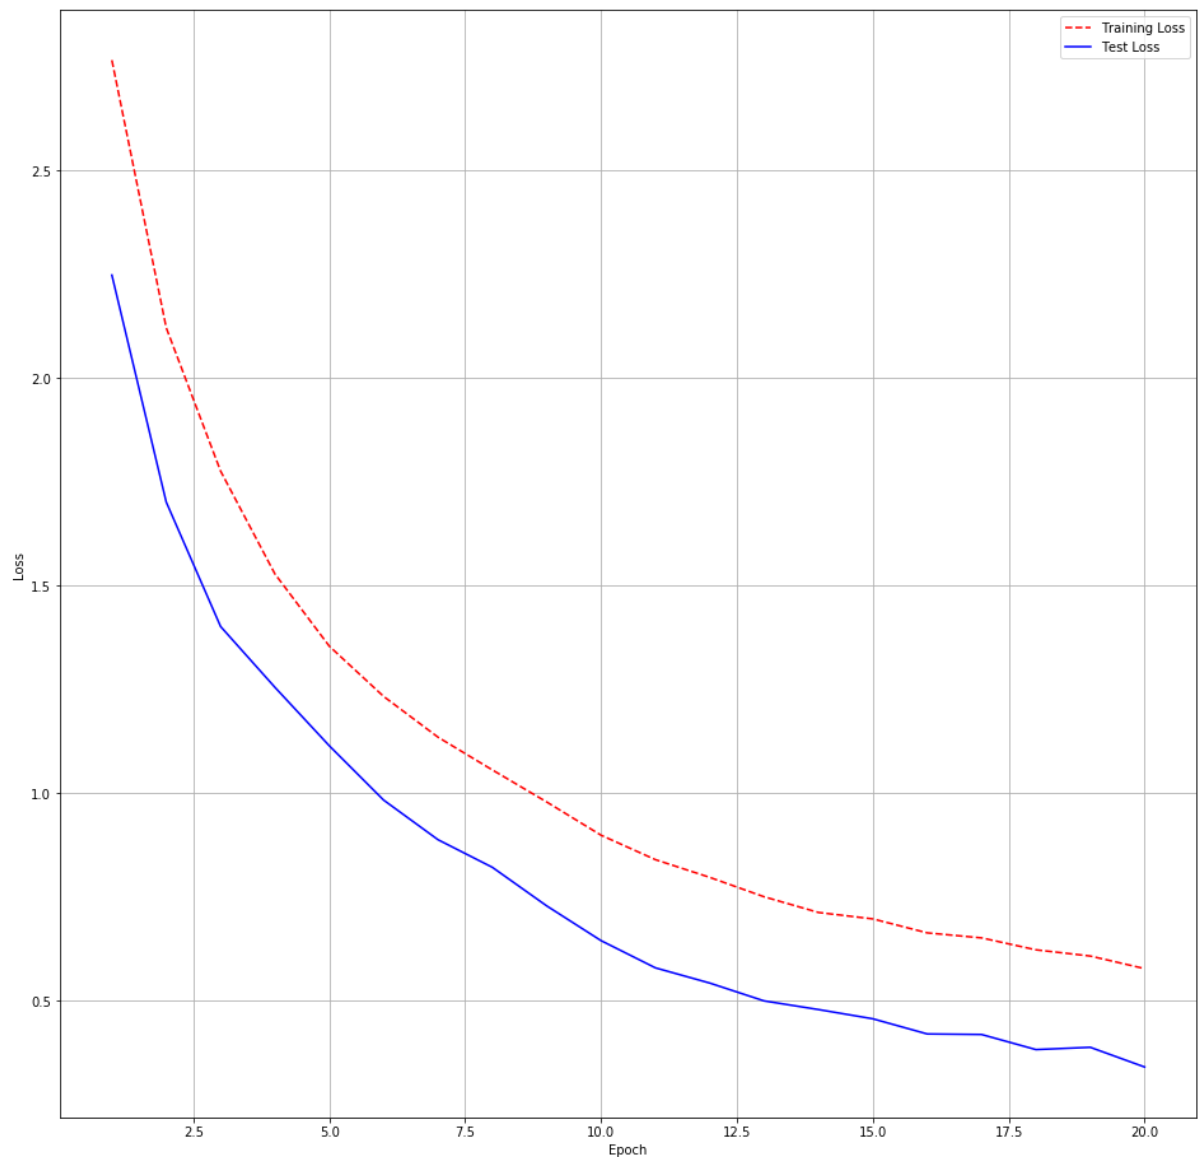
- acc: 0.6512 - val\_loss: 0.7290 - val\_acc: 0.7849  
Epoch 10/20  
60000/60000 [=====] - 4s 74us/step - loss: 0.8991  
- acc: 0.6950 - val\_loss: 0.6454 - val\_acc: 0.8035  
Epoch 11/20  
60000/60000 [=====] - 4s 72us/step - loss: 0.8403  
- acc: 0.7196 - val\_loss: 0.5799 - val\_acc: 0.8399  
Epoch 12/20  
60000/60000 [=====] - 4s 73us/step - loss: 0.7977  
- acc: 0.7374 - val\_loss: 0.5432 - val\_acc: 0.8696  
Epoch 13/20  
60000/60000 [=====] - 4s 69us/step - loss: 0.7509  
- acc: 0.7566 - val\_loss: 0.5003 - val\_acc: 0.8430  
Epoch 14/20  
60000/60000 [=====] - 4s 75us/step - loss: 0.7132  
- acc: 0.7731 - val\_loss: 0.4795 - val\_acc: 0.8792  
Epoch 15/20  
60000/60000 [=====] - 4s 69us/step - loss: 0.6977  
- acc: 0.7852 - val\_loss: 0.4573 - val\_acc: 0.8894  
Epoch 16/20  
60000/60000 [=====] - 4s 71us/step - loss: 0.6641  
- acc: 0.8003 - val\_loss: 0.4208 - val\_acc: 0.9031  
Epoch 17/20  
60000/60000 [=====] - 4s 72us/step - loss: 0.6520  
- acc: 0.8060 - val\_loss: 0.4195 - val\_acc: 0.8997  
Epoch 18/20  
60000/60000 [=====] - 4s 74us/step - loss: 0.6233  
- acc: 0.8197 - val\_loss: 0.3831 - val\_acc: 0.9206  
Epoch 19/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.6084  
- acc: 0.8303 - val\_loss: 0.3886 - val\_acc: 0.9292  
Epoch 20/20  
60000/60000 [=====] - 4s 72us/step - loss: 0.5780  
- acc: 0.8404 - val\_loss: 0.3412 - val\_acc: 0.9306

```
In [30]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_42.history['loss']
test_loss = history_42.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_42.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.3411893347978592

Test accuracy: 0.9306

## **[4.3] 3-Layer MLP + ReLu + Adam + Dropout + Batch Normalization**

```
In [31]: mdl_relu_43 = Sequential()
mdl_relu_43.add(Dense(512, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_43.add(Dropout(rate=0.5))
mdl_relu_43.add(BatchNormalization())
mdl_relu_43.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_43.add(Dropout(rate=0.5))
mdl_relu_43.add(BatchNormalization())
mdl_relu_43.add(Dense(64, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.175, seed=42)))
mdl_relu_43.add(Dropout(rate=0.5))
mdl_relu_43.add(BatchNormalization())
mdl_relu_43.add(Dense(32, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.180, seed=42)))
mdl_relu_43.add(Dropout(rate=0.5))
mdl_relu_43.add(BatchNormalization())
mdl_relu_43.add(Dense(16, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.190, seed=42)))
mdl_relu_43.add(Dropout(rate=0.5))
mdl_relu_43.add(BatchNormalization())
mdl_relu_43.add(Dense(output_dim,activation='softmax'))
mdl_relu_43.summary()
mdl_relu_43.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_43 = mdl_relu_43.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch,verbose=1,validation_data=(x_test,y_test))
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
dense_29 (Dense)	(None, 512)	401920
dropout_12 (Dropout)	(None, 512)	0
batch_normalization_7 (Batch Normalization)	(None, 512)	2048
dense_30 (Dense)	(None, 128)	65664
dropout_13 (Dropout)	(None, 128)	0
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dense_31 (Dense)	(None, 64)	8256
dropout_14 (Dropout)	(None, 64)	0
batch_normalization_9 (Batch Normalization)	(None, 64)	256
dense_32 (Dense)	(None, 32)	2080
dropout_15 (Dropout)	(None, 32)	0
batch_normalization_10 (Batch Normalization)	(None, 32)	128
dense_33 (Dense)	(None, 16)	528
dropout_16 (Dropout)	(None, 16)	0
batch_normalization_11 (Batch Normalization)	(None, 16)	64
dense_34 (Dense)	(None, 10)	170
Total params: 481,626		
Trainable params: 480,122		
Non-trainable params: 1,504		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 193us/step - loss: 2.148

7 - acc: 0.2260 - val\_loss: 1.3847 - val\_acc: 0.6565

Epoch 2/20

60000/60000 [=====] - 9s 156us/step - loss: 1.5526

- acc: 0.4431 - val\_loss: 0.9069 - val\_acc: 0.7928

Epoch 3/20

60000/60000 [=====] - 9s 154us/step - loss: 1.2466

- acc: 0.5564 - val\_loss: 0.6729 - val\_acc: 0.8246

Epoch 4/20

60000/60000 [=====] - 9s 157us/step - loss: 1.0757

- acc: 0.6197 - val\_loss: 0.5228 - val\_acc: 0.8483

Epoch 5/20

60000/60000 [=====] - 10s 168us/step - loss: 0.970

8 - acc: 0.6539 - val\_loss: 0.4565 - val\_acc: 0.8560

Epoch 6/20



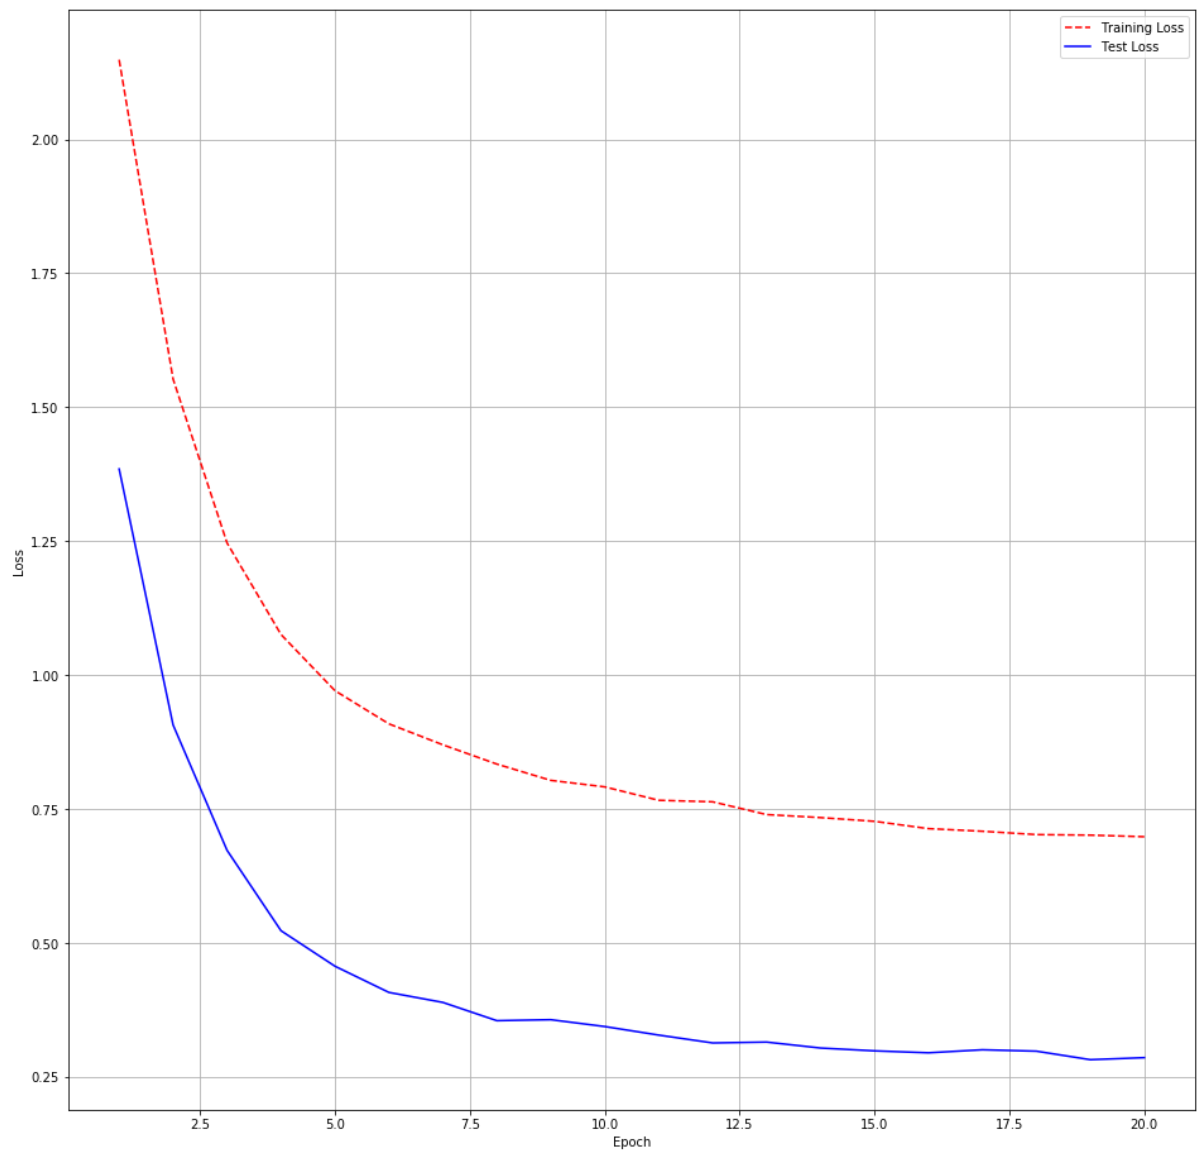
```
60000/60000 [=====] - 9s 156us/step - loss: 0.9089
- acc: 0.6775 - val_loss: 0.4076 - val_acc: 0.8588
Epoch 7/20
60000/60000 [=====] - 10s 163us/step - loss: 0.869
8 - acc: 0.6878 - val_loss: 0.3889 - val_acc: 0.8634
Epoch 8/20
60000/60000 [=====] - 9s 156us/step - loss: 0.8338
- acc: 0.6996 - val_loss: 0.3548 - val_acc: 0.8719
Epoch 9/20
60000/60000 [=====] - 10s 160us/step - loss: 0.803
4 - acc: 0.7104 - val_loss: 0.3566 - val_acc: 0.8721
Epoch 10/20
60000/60000 [=====] - 9s 151us/step - loss: 0.7913
- acc: 0.7120 - val_loss: 0.3439 - val_acc: 0.8742
Epoch 11/20
60000/60000 [=====] - 9s 156us/step - loss: 0.7662
- acc: 0.7214 - val_loss: 0.3279 - val_acc: 0.8767
Epoch 12/20
60000/60000 [=====] - 9s 155us/step - loss: 0.7634
- acc: 0.7200 - val_loss: 0.3133 - val_acc: 0.8794
Epoch 13/20
60000/60000 [=====] - 9s 150us/step - loss: 0.7395
- acc: 0.7262 - val_loss: 0.3148 - val_acc: 0.8792
Epoch 14/20
60000/60000 [=====] - 10s 164us/step - loss: 0.733
9 - acc: 0.7276 - val_loss: 0.3037 - val_acc: 0.8816
Epoch 15/20
60000/60000 [=====] - 10s 162us/step - loss: 0.727
0 - acc: 0.7300 - val_loss: 0.2985 - val_acc: 0.8827
Epoch 16/20
60000/60000 [=====] - 9s 154us/step - loss: 0.7133
- acc: 0.7360 - val_loss: 0.2949 - val_acc: 0.8851
Epoch 17/20
60000/60000 [=====] - 9s 154us/step - loss: 0.7084
- acc: 0.7366 - val_loss: 0.3005 - val_acc: 0.8811
Epoch 18/20
60000/60000 [=====] - 9s 155us/step - loss: 0.7022
- acc: 0.7363 - val_loss: 0.2979 - val_acc: 0.8818
Epoch 19/20
60000/60000 [=====] - 9s 154us/step - loss: 0.7012
- acc: 0.7376 - val_loss: 0.2820 - val_acc: 0.8871
Epoch 20/20
60000/60000 [=====] - 10s 164us/step - loss: 0.698
2 - acc: 0.7382 - val_loss: 0.2858 - val_acc: 0.8852
```

```
In [32]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_43.history['loss']
test_loss = history_43.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_43.evaluate(x_test, y_test, verbose=0)
plt_mdl_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.28580667269229887

Test accuracy: 0.8852

## **[4.4] 3-Layer MLP + ReLu + Adam + Batch Normalization**

**No Dropout**

```
In [33]: mdl_relu_44 = Sequential()
mdl_relu_44.add(Dense(512, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(mean=0.0,stddev=0.125, seed=42)))
mdl_relu_44.add(BatchNormalization())
mdl_relu_44.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.150, seed=42)))
mdl_relu_44.add(BatchNormalization())
mdl_relu_44.add(Dense(64, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.175, seed=42)))
mdl_relu_44.add(BatchNormalization())
mdl_relu_44.add(Dense(32, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.180, seed=42)))
mdl_relu_44.add(BatchNormalization())
mdl_relu_44.add(Dense(16, activation='relu',kernel_initializer=RandomNormal(mean=0.0,stddev=0.190, seed=42)))
mdl_relu_44.add(BatchNormalization())
mdl_relu_44.add(Dense(output_dim,activation='softmax'))
mdl_relu_44.summary()
mdl_relu_44.compile(optimizer='adam',loss='categorical_crossentropy', metrics=['accuracy'])
history_44 = mdl_relu_44.fit(X_train,y_train,batch_size=batch_size,epochs=no_epoch,verbose=1,validation_data=(x_test,y_test))
```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
batch_normalization_12 (Batch Normalization)	(None, 512)	2048
dense_36 (Dense)	(None, 128)	65664
batch_normalization_13 (Batch Normalization)	(None, 128)	512
dense_37 (Dense)	(None, 64)	8256
batch_normalization_14 (Batch Normalization)	(None, 64)	256
dense_38 (Dense)	(None, 32)	2080
batch_normalization_15 (Batch Normalization)	(None, 32)	128
dense_39 (Dense)	(None, 16)	528
batch_normalization_16 (Batch Normalization)	(None, 16)	64
dense_40 (Dense)	(None, 10)	170

Total params: 481,626  
 Trainable params: 480,122  
 Non-trainable params: 1,504

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 189us/step - loss: 0.4369 - acc: 0.8869 - val\_loss: 0.1484 - val\_acc: 0.9578

Epoch 2/20

60000/60000 [=====] - 9s 155us/step - loss: 0.1202 - acc: 0.9670 - val\_loss: 0.1058 - val\_acc: 0.9687

Epoch 3/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0718 - acc: 0.9797 - val\_loss: 0.0920 - val\_acc: 0.9715

Epoch 4/20

60000/60000 [=====] - 9s 149us/step - loss: 0.0522 - acc: 0.9845 - val\_loss: 0.0946 - val\_acc: 0.9715

Epoch 5/20

60000/60000 [=====] - 9s 149us/step - loss: 0.0393 - acc: 0.9884 - val\_loss: 0.0861 - val\_acc: 0.9756

Epoch 6/20

60000/60000 [=====] - 9s 149us/step - loss: 0.0335 - acc: 0.9897 - val\_loss: 0.0817 - val\_acc: 0.9768

Epoch 7/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0276 - acc: 0.9914 - val\_loss: 0.0832 - val\_acc: 0.9762

Epoch 8/20

60000/60000 [=====] - 9s 152us/step - loss: 0.0233 - acc: 0.9924 - val\_loss: 0.0878 - val\_acc: 0.9743

Epoch 9/20

60000/60000 [=====] - 9s 156us/step - loss: 0.0252

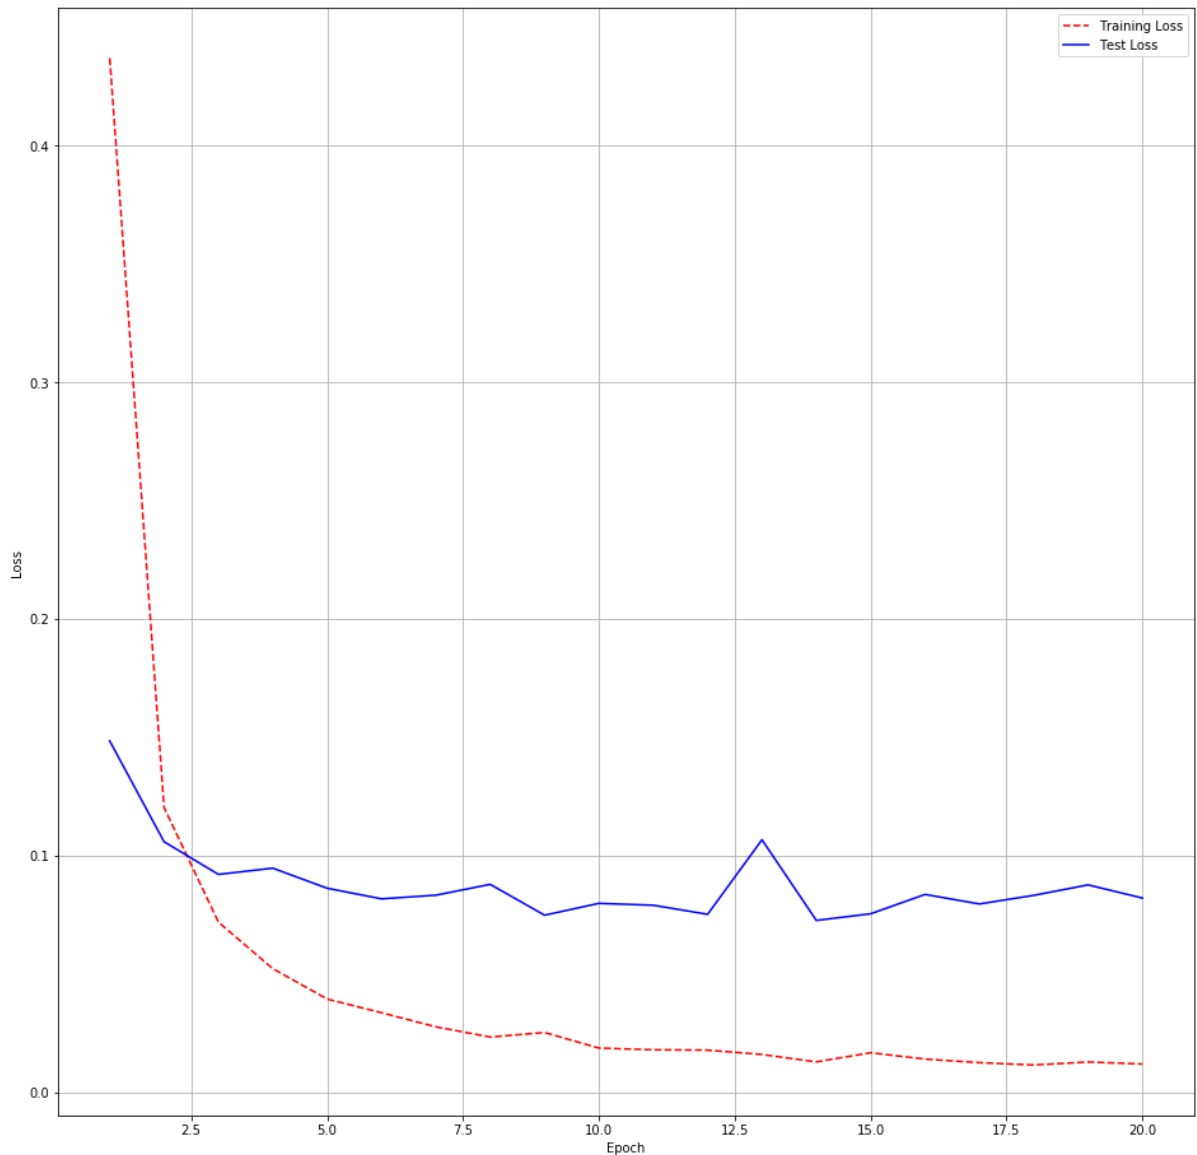
```
- acc: 0.9920 - val_loss: 0.0747 - val_acc: 0.9801
Epoch 10/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0186
- acc: 0.9938 - val_loss: 0.0797 - val_acc: 0.9785
Epoch 11/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0179
- acc: 0.9940 - val_loss: 0.0790 - val_acc: 0.9796
Epoch 12/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0177
- acc: 0.9939 - val_loss: 0.0751 - val_acc: 0.9825
Epoch 13/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0159
- acc: 0.9948 - val_loss: 0.1066 - val_acc: 0.9744
Epoch 14/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0128
- acc: 0.9960 - val_loss: 0.0725 - val_acc: 0.9818
Epoch 15/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0166
- acc: 0.9945 - val_loss: 0.0753 - val_acc: 0.9797
Epoch 16/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0139
- acc: 0.9955 - val_loss: 0.0835 - val_acc: 0.9809
Epoch 17/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0124
- acc: 0.9960 - val_loss: 0.0795 - val_acc: 0.9796
Epoch 18/20
60000/60000 [=====] - 10s 159us/step - loss: 0.011
4 - acc: 0.9962 - val_loss: 0.0831 - val_acc: 0.9806
Epoch 19/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0127
- acc: 0.9959 - val_loss: 0.0875 - val_acc: 0.9792
Epoch 20/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0119
- acc: 0.9961 - val_loss: 0.0820 - val_acc: 0.9800
```

```
In [34]: %matplotlib inline
import keras
from matplotlib import pyplot as plt

epoch_count = list(range(1,no_epoch+1))

# Get training and test loss histories
training_loss = history_44.history['loss']
test_loss = history_44.history['val_loss']

# Create count of the number of epochs
score = mdl_relu_44.evaluate(x_test, y_test, verbose=0)
plt_md1_res(epoch_count, training_loss, test_loss, score[0], score[1])
```



Test score: 0.08198474683784589

Test accuracy: 0.98

## Conclusion

```
In [1]: import tabulate
```

## 2- Hidden Layer results

```
In [9]: res_no_2nn =[['Hidden \nLayer', 'Number of \nNeurons'],  
                    [1,512],  
                    [2,128]  
                    ]  
print(tabulate.tabulate(res_no_2nn, tablefmt='fancy_grid'))
```

Hidden Layer	Number of Neurons
1	512
2	128

```
In [10]: res_tab_2 =[['Dropout','Batch \nNormalization', 'Test \nAccuracy'],  
                    ['No','No',0.9796],  
                    ['Yes','No',0.9816],  
                    ['No','Yes',0.981],  
                    ['Yes','Yes',0.9809]  
                    ]  
print(tabulate.tabulate(res_tab_2,tablefmt='fancy_grid'))
```

Dropout	Batch Normalization	Test Accuracy
No	No	0.9796
Yes	No	0.9816
No	Yes	0.981
Yes	Yes	0.9809

## 3- Hidden Layer results



```
In [7]: res_no_3nn =[['Hidden \nLayer', 'Number of \nNeurons'],  
                    [1,1024],  
                    [2,512],  
                    [3,128]  
                    ]  
print(tabulate.tabulate(res_no_3nn, tablefmt='fancy_grid'))
```

Hidden Layer	Number of Neurons
1	1024
2	512
3	128

```
In [8]: res_tab_3 =[['Dropout','Batch \nNormalization', 'Test \nAccuracy'],  
                   ['No','No',0.9798],  
                   ['Yes','No',0.9797],  
                   ['No','Yes',0.9771],  
                   ['Yes','Yes',0.9846]  
                   ]  
print(tabulate.tabulate(res_tab_3,tablefmt='fancy_grid'))
```

Dropout	Batch Normalization	Test Accuracy
No	No	0.9798
Yes	No	0.9797
No	Yes	0.9771
Yes	Yes	0.9846

## 5- Hidden Layer results

```
In [5]: res_no_5nn =[['Hidden \nLayer', 'Number of \nNeurons'],
                    [1,512],
                    [2,128],
                    [3,64],
                    [4,32],
                    [5,16]
                    ]
print(tabulate.tabulate(res_no_5nn, tablefmt='fancy_grid'))
```

Hidden Layer	Number of Neurons
1	512
2	128
3	64
4	32
5	16

```
In [6]: res_tab_5 =[['Dropout','Batch \nNormalization', 'Test \nAccuracy'],
                    ['No','No',0.9872],
                    ['Yes','No',0.9306],
                    ['No','Yes',0.98],
                    ['Yes','Yes',0.8852]
                    ]
print(tabulate.tabulate(res_tab_5,tablefmt='fancy_grid'))
```

Dropout	Batch Normalization	Test Accuracy
No	No	0.9872
Yes	No	0.9306
No	Yes	0.98
Yes	Yes	0.8852

**The best and worst performance comes from the model that has 5-hidden layers.**

**The best performance model is the one that has no Dropout or Batch Normalization with a Test-accuracy of 0.9872**

**The worst performance model has both Dropout and Batch Normalization enabled with a Test-accuracy of 0.8852**